



#codemash

MongoDB – Zero to Sharding

Sridhar Nanjundeswaran

Engineer, MongoDB Inc.

@snanjund

Workshop Agenda

- Introduction and building your first app
- Schema Design and Indexing
- Replication and Sharding
- Deployment ?

Part 1 – Introduction and building your first app

This Talk

- Quick introduction to MongoDB
- Data modeling in MongoDB, queries, geospatial, updates and map reduce.
- Using a location-based app as an example
- https://github.com/sridharn/codemash_2014/tree/master/firstapp

Relational Database Challenges

Data Types

- Unstructured data
- Semi-structured data
- Polymorphic data

Volume of Data

- Petabytes of data
- Trillions of records
- Tens of millions of queries per second



Agile Development

- Iterative
- Short development cycles
- New workloads

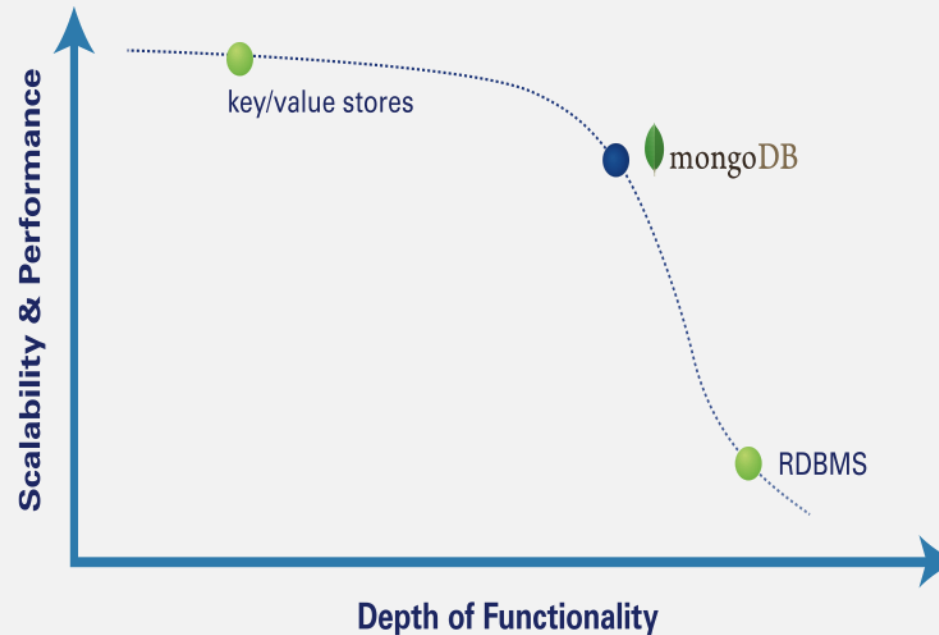
New Architectures

- Horizontal scaling
- Commodity servers
- Cloud computing

What is mongoDB?

MongoDB is a scalable, high-performance, open source, document database.

- Fast Querying
- In-place updates
- Full Index Support
- Replication /High Availability
- Auto-Sharding
- Aggregation; Map/Reduce
- DBMS of the year 2013 per DB-Engines.com 😊



Document database?

- JSON - json.org
 - { “name” : “MongoDB” }
- BSON - bsonspec.org
 - “\x16\x00\x00\x00\x02_id\x00\x08\x00\x00\x00MongoDB\x00\x00”
- BSON
 - Storage format
 - In wire protocol

Why BSON?

- JSON has powerful, limited set of datatypes
 - Mongo extends datatypes with Date, Int types, ObjectId,
- BSON is a binary representation of JSON
 - Optimized for performance and navigational abilities
 - Also compression

Where can you use it?


- MongoDB is Implemented in C++
- Windows, Linux, Mac OS-X, Solaris



- Packages available
 - OS X – Macports, Homebrew
 - Linux – Debian, Ubuntu, Fedora, CentOS...
 - Windows – MSIs (coming soon)

How can I connect to it?

Official MongoDB drivers



MongoDB Drivers

- Official Support for 10 languages
- Community drivers for tons more
 - R, lua etc.
- Drivers connect to mongo servers
- Drivers translate BSON into native types
- mongo shell is not a driver, but works like one in some ways
- Installed using typical means (npm, pecl, gem, pip)

Terminology

RDBMS	MongoDB
Table/View	Collection
Row(s)	Document
Index	Index
Partition	Shard
Join	Embedding/Linking
Fixed Schema	Flexible/Implied Schema

Example Document

```
{
  _id : ObjectId("4c4ba5c0672c685e5e8aabf3"),
  author : "Sridhar",
  date : ISODate("2014-01-02T11:52:27.442Z"),
  text : "About MongoDB...",
  tags : [ "tech", "databases", "nosql" ],
  comments : [{
    author : "Doug",
    date : ISODate("2014-01-03T17:22:21.124Z"), text
: "Best Post Ever!"
  }],
  comment_count : 1
}
```

Why use MongoDB?

- Intrinsic support for fast, iterative development
- Super low latency access to your data
- Very little CPU overhead
- No additional caching layer required
- Built in replication and horizontal scaling support

Building your first app

Install MongoDB

1. Download
2. Unzip
3. Create data directory
4. Run mongod

Sample scripts at

https://github.com/sridharn/codemash_2014/tree/master/firstapp

Yes it is as simple as that 😊

Building Your First MongoDB App

- Want to build an app where users can check in to a location



- Leave notes or comments about that location

Requirements

- "As a user I want to be able to find other locations nearby"
- Need to store locations (Offices, Restaurants, etc)
 - name, address, tags
 - coordinates
 - User generated content e.g. tips / notes

Requirements

"As a user I want to be able to 'checkin' to a location"

Checkins

- User should be able to 'check in' to a location
- Want to be able to generate statistics:
 - Recent checkins
 - Popular locations

Collections

loc1, loc2, loc3

locations

user1, user2

users

checkin1, checkin2

checkins

Locations v1

```
> location_1 = {  
    name: "Taj Mahal",  
    address: "123 University Ave",  
    city: "Palo Alto",  
    zipcode: 94301  
}
```

Locations v1

```
> location_1 = {  
    name: "Taj Mahal",  
    address: "123 University Ave",  
    city: "Palo Alto",  
    zipcode: 94301  
}  
  
> db.locations.insert(location_1)
```

Locations v1

```
> db.locations.findOne({name: "Taj Mahal"})
```


```
{  
  "_id" : ObjectId("50e67a4f4b23019a4ab9b58c"),  
  "name" : "Taj Mahal",  
  "address" : "123 University Ave",  
  "city" : "Palo Alto",  
  "zipcode" : 94301  
}
```

What is _id?

- _id is the primary key in MongoDB
- Automatically indexed
- Automatically created as an ObjectId if not provided
- Any unique immutable value could be used

What is ObjectId?

- ObjectId is a special 12 byte value
- Guaranteed to be unique across your cluster
- `ObjectId("50e67a4f4b23019a4ab9b58c")`



Locations v1 – Indexed find

```
> db.locations.ensureIndex({name: 1})
```

```
> db.locations.find({name: "Taj  
Mahal"}).explain()
```

```
{
```

```
  "cursor" : "BtreeCursor name_1",
```

```
  "isMultiKey" : false,
```

```
  ...
```

Locations v2

```
> location_2 = {  
    name: "Lotus Flower",  
    address: "234 University Ave",  
    city: "Palo Alto",  
    zipcode: 94301,  
    tags: ["restaurant", "dumplings"]  
}  
  
> db.locations.insert(location_2)
```

Locations v2

```
> db.locations.findOne({tags: "dumplings"})
{
  "_id" : ObjectId("50e67f334b23019a4ab9b59a"),
  "name" : "Lotus Flower",
  ...
}

> db.locations.ensureIndex({tags: 1})

> db.locations.find({tags: "dumplings"}).explain()
{
  "cursor" : "BtreeCursor tags_1",
  "isMultiKey" : true,
  ...
}
```

Locations v3

```
> location_3 = {  
    name: "El Capitan",  
    address: "345 University Ave",  
    city: "Palo Alto",  
    zipcode: 94301,  
    tags: ["restaurant", "tacos"],  
    lat_long: [52.5184, 13.387]  
}  
  
> db.locations.insert(location_3)
```

Locations v3

```
> db.locations.find({lat_long: {$near:[52.53, 13.4]}})

error: {
  "$err" : "can't find special index: 2d for: {
lat_long: { $near: [ 52.53, 13.4 ] } }",
  "code" : 13038
}

> db.locations.ensureIndex({lat_long: "2d"})

> db.locations.findOne({lat_long: {$near:[52.53,
13.4]}})

{
  "_id" : ObjectId("50e686ab4b23019a4ab9b59d"),
  "name" : "El Capitan",
  ...
```

Finding locations

// creating your indexes:

```
> db.locations.ensureIndex({tags: 1})
```

```
> db.locations.ensureIndex({name: 1})
```

```
> db.locations.ensureIndex({lat_long: "2d"})
```

// finding places:

```
> db.locations.find({lat_long: {$near: [52.53, 13.4]}})
```

// with regular expressions:

```
> db.locations.find({name: /^Taj/})
```

// by tag:

```
> db.locations.find({tag: "dumplings"})
```

Updating Documents

Atomic operators:

\$set, \$unset, \$inc, \$push, \$pushAll, \$pull, \$pullAll

Locations - adding tips

// adding a tip with update:

```
> db.locations.update(  
  {name: "Lotus Flower"},  
  {$push: {  
    tips: {  
      user: "Sridhar",  
      date: ISODate("2012-09-21T11:52:27.442Z"),  
      tip: "The sesame dumplings are  
awesome!" }  
    } })
```

task - done

```
> db.locations.findOne({name:/^Lot/})
{
  "_id" : ObjectId("50e67f334b23019a4ab9b59a"),
  "address" : "234 University Ave",
  "city" : "Palo Alto",
  "name" : "Lotus Flower",
  "tags" : [
    "restaurant",
    "dumplings"
  ],
  "tips" : [
    {
      "user" : "Sridhar",
      "date" : ISODate("2012-09-21T11:52:27.442Z"),
      "tip" : "The sesame dumplings are awesome!"
    }
  ],
  "zipcode" : 94301
}
```



Requirements

"As a user I want to be able to 'checkin' to a location"

Checkins

- User should be able to 'check in' to a location
- Want to be able to generate statistics:
 - Recent checkins
 - Popular locations

Users and Checkins

```
> user_1 = {  
  _id: "sridhar@10gen.com",  
  name: "Sridhar",  
  twitter: "snanjund",  
  checkins: [  
    {location: "Lotus Flower", ts: ISODate("2012-09-  
21T11:52:27.442Z")},  
    {location: "Taj Mahal", ts: ISODate("2012-09-  
22T07:15:00.442Z")}  
  ]  
}  
  
> db.users.save(user_1)  
  
> db.users.ensureIndex({"checkins.location": 1})
```

Simple Stats

```
// find all users who've checked in here:
```

```
> db.users.find({"checkins.location":"Lotus Flower"},  
  {name:1, checkins:1})
```

```
// find the last 10 checkins here:
```

```
> db.users.find({"checkins.location":"Lotus Flower"},  
  {name:1, checkins:1}).sort({"checkins.ts": -  
  1}).limit(10)
```

Hard to query for last 10

User and Checkins v2

```
> user_1 = {
    _id: "sridhar@10gen.com",
    name: "Sridhar",
    twitter: "snanjund",
}

> location_id = db.locations.findOne({name: "Taj
Mahal"}, {_id:1})["_id"]

> checkin_1 = {
    location: location_id,
    user: "sridhar@10gen.com",
    ts: ISODate("2012-09-21T11:52:27.442Z")
}
```

Simple Stats

// find all users who've checked in here:

```
> location_id = db.locations.find({"name": "Lotus Flower"})
> u_ids = db.checkins.find({location: location_id,
                           {_id: -1, user: 1}})

> users = db.users.find({_id: {$in: u_ids}})
```

// find the last 10 checkins here:

```
> db.checkins.find({location: location_id})
                           .sort({ts: -1}).limit(10)
```

// count how many checked in today:

```
> db.checkins.find({location: location_id,
                   ts: {$gt: midnight}}
                  ).count()
```

Aggregation- in Mongo 2.2+

// Find most popular locations

```
> agg = db.checkins.aggregate(  
    {$match: {ts: {$gt: now_minus_3_hrs}}},  
    {$group: {_id: "$location", numEntries:  
{$sum: 1}}}  
)  
  
> agg.result  
[{"_id": "Lotus Flower", "numEntries" : 17}]
```


Map Reduce

```
// Find most popular locations
```

```
> map_func = function() {  
    emit(this.location, 1);  
}  
  
> reduce_func = function(key, values) {  
    return Array.sum(values);  
}  
  
> db.checkins.mapReduce(map_func, reduce_func,  
    {query: {ts: {$gt: now_minus_3_hrs}},  
    out: "result"})  
  
> db.result.findOne()  
{"_id": "Lotus Flower", "value" : 17}
```



#codemash

Coming Next:

Part 2 : Schema Design and Indexing

Sridhar Nanjundeswaran

Engineer, MongoDB Inc.

@snanjund