



#codemash

MongoDB – Zero to Sharding

Part 3.1 : Replication

Sridhar Nanjundeswaran

Engineer, MongoDB Inc.

@snanjund

Agenda

- Introduction to Replica Sets
- Developing with Replica Sets
- Operational Considerations
- https://github.com/sridharn/codemash_2014/tree/master/replication




Introduction to Replica Sets

- Why?
- What is it?
- Configuration Options


Why Replication?

- How many have faced node failures?
- How many have been woken to do fail overs?
- How many have experienced issues due to n/w latency?
- Different uses for data
 - Normal processing
 - Simple analytics


Replica Set - Creation



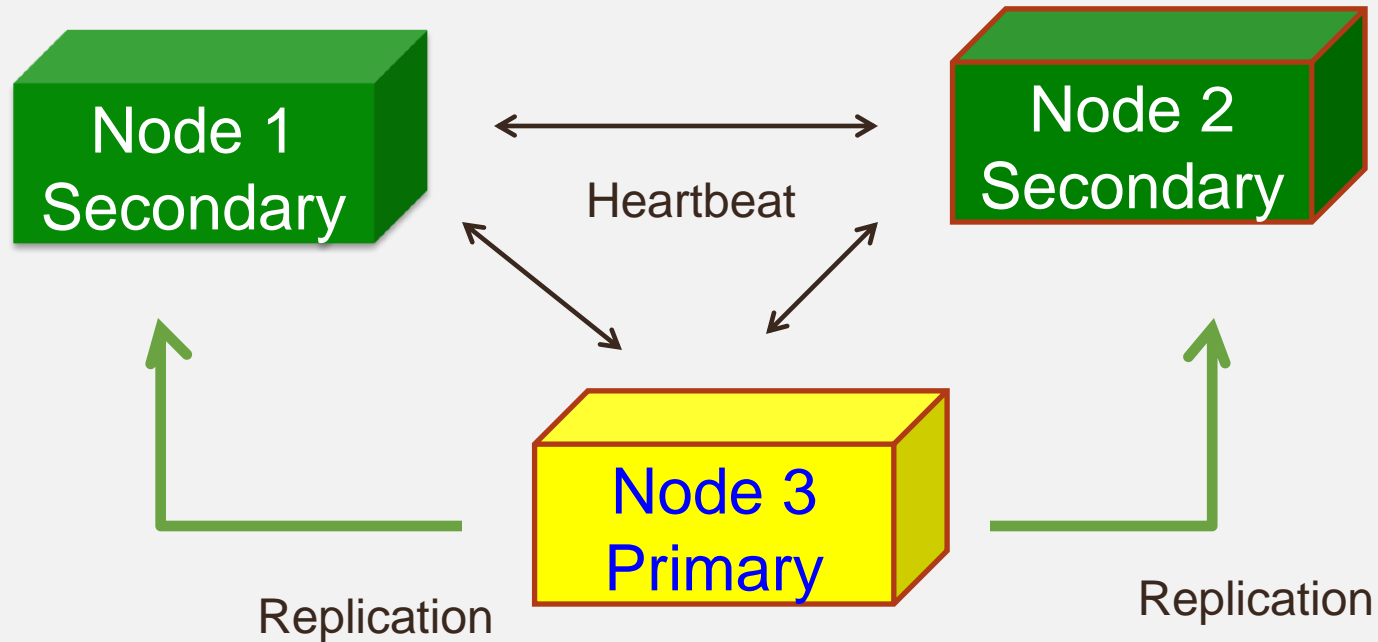
Replica Set - Initialize



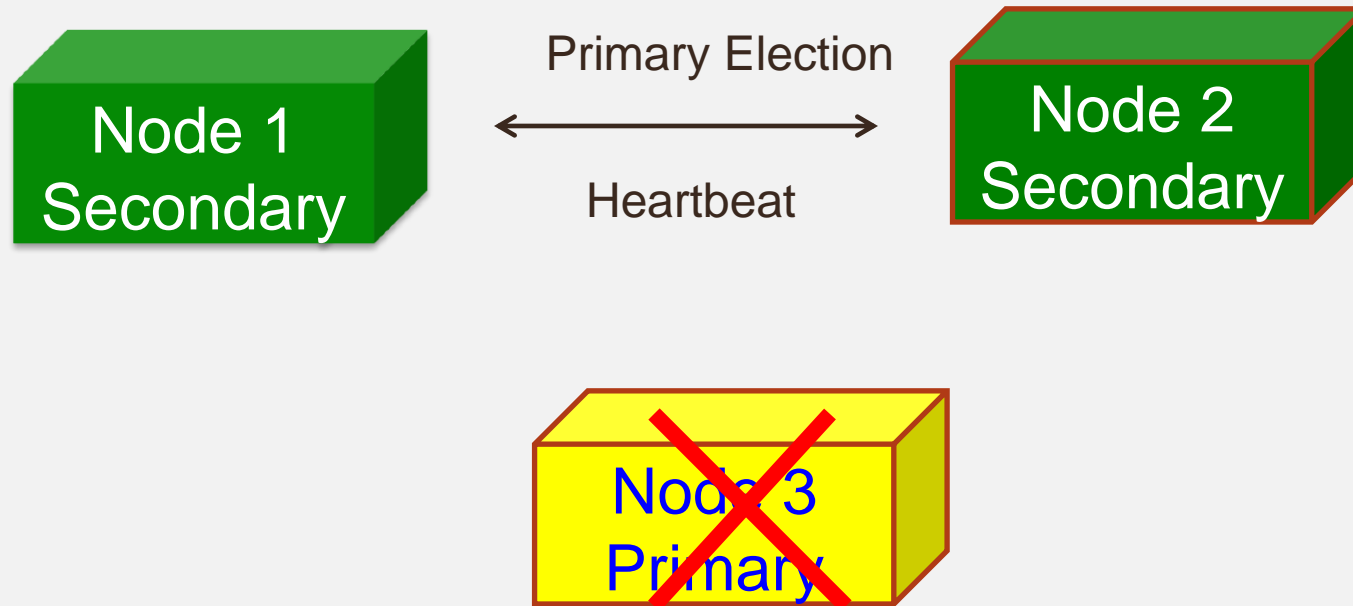
Replica Set - Initializing



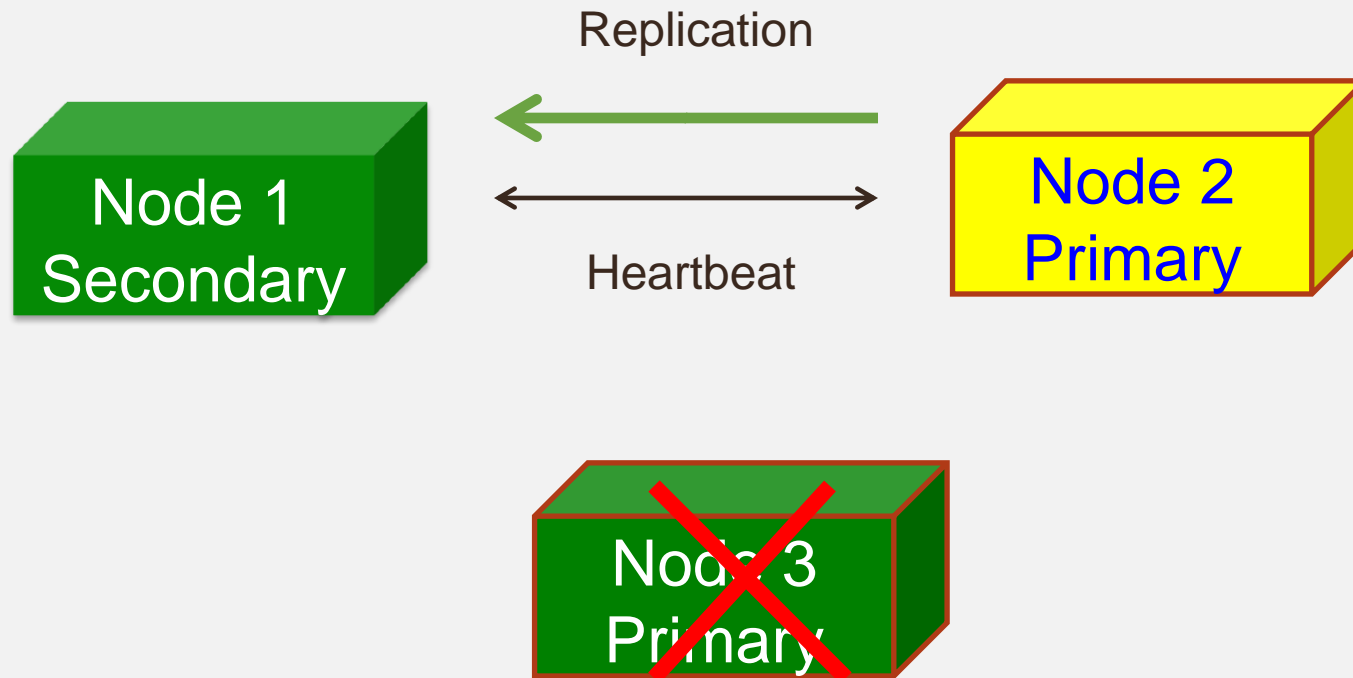
Replica Set - Initialized




Replica Set - Failure



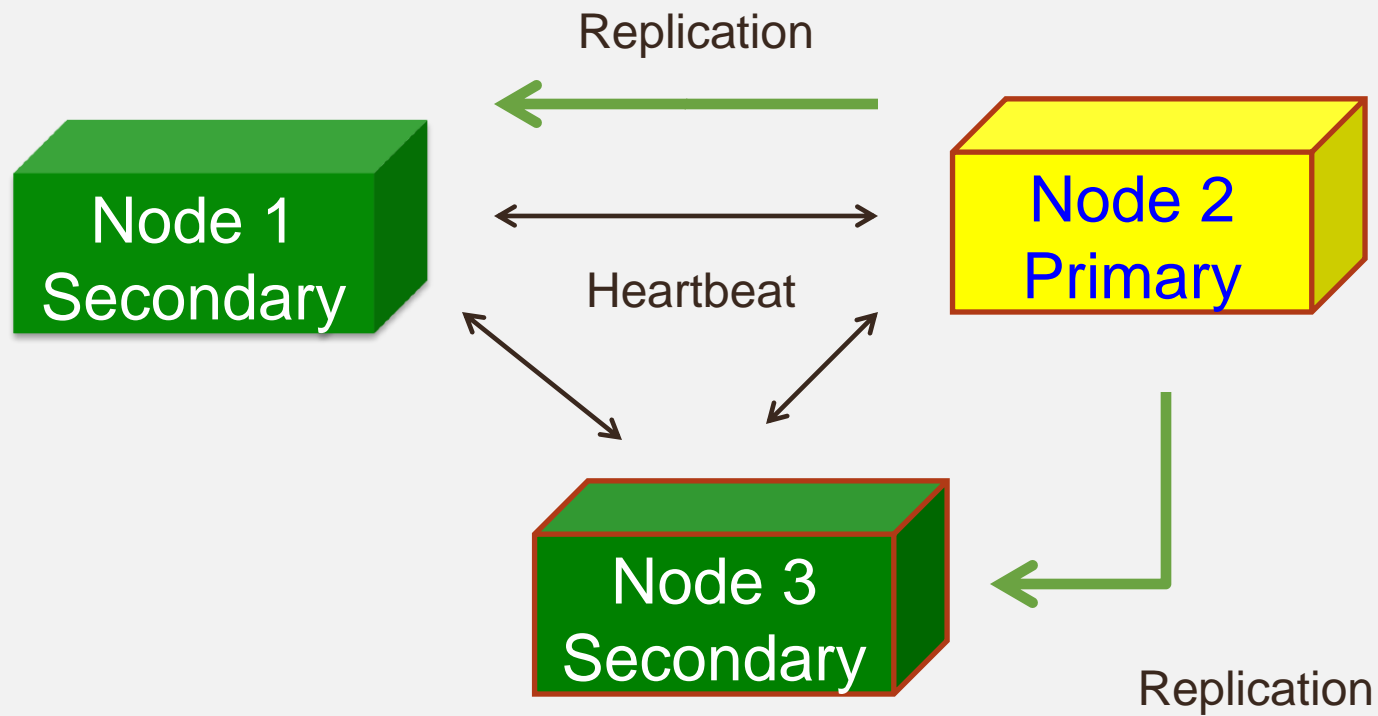
Replica Set - Failover



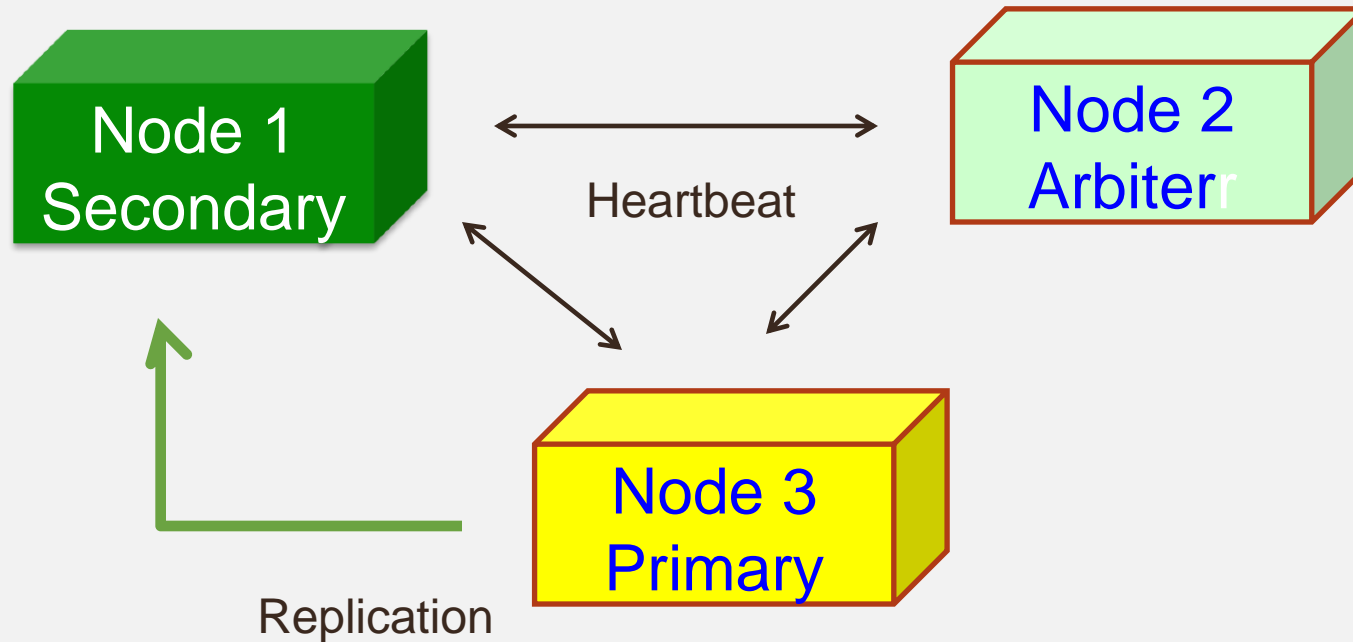
Replica Set - Recovery



Replica Set – Recovery Complete



Replica Set – Member Roles



Replica Set – Configuration Options

```
> conf = {  
  _id : "mySet",  
  members : [  
    { _id : 0, host : "A", priority : 3 },  
    { _id : 1, host : "B", priority : 2 },  
    { _id : 2, host : "C" },  
    { _id : 3, host : "D", hidden : true },  
    { _id : 4, host : "E", hidden : true, slaveDelay : 3600 }  
  ]  
}  
  
> rs.initiate(conf)
```

Replica Set – Configuration Options

```
> conf = {  
  _id : "mySet",  
  members : [  
    { _id : 0, host : "A", priority : 3 },  
    { _id : 1, host : "B", priority : 2 },  
    { _id : 2, host : "C" },  
    { _id : 3, host : "D", hidden : true },  
    { _id : 4, host : "E", hidden : true, slaveDelay : 3600 }  
  ]  
}
```

Primary DC



```
> rs.initiate(conf)
```

Replica Set – Configuration Options

```
> conf = {  
  _id : "mySet",  
  members : [  
    { _id : 0, host : "A", priority : 3 },  
    { _id : 1, host : "B", priority : 2 },  
    { _id : 2, host : "C" },  
    { _id : 3, host : "D", hidden : true },  
    { _id : 4, host : "E", hidden : true, slaveDelay : 3600 }  
  ]  
}
```

**Secondary DC
Default priority = 1**



```
> rs.initiate(conf)
```


Replica Set – Configuration Options

```
> conf = {  
  _id : "mySet",  
  members : [  
    { _id : 0, host : "A", priority : 3 },  
    { _id : 1, host : "B", priority : 2 },  
    { _id : 2, host : "C" },  
    { _id : 3, host : "D", hidden : true },  
    { _id : 4, host : "E", hidden : true, slaveDelay : 3600 }  
  ]  
}
```

Analytics node



```
> rs.initiate(conf)
```

Replica Set – Configuration Options

```
> conf = {  
  _id : "mySet",  
  members : [  
    { _id : 0, host : "A", priority : 3 },  
    { _id : 1, host : "B", priority : 2 },  
    { _id : 2, host : "C" },  
    { _id : 3, host : "D", hidden : true },  
    { _id : 4, host : "E", hidden : true, slaveDelay : 3600 }  
  ]  
}
```

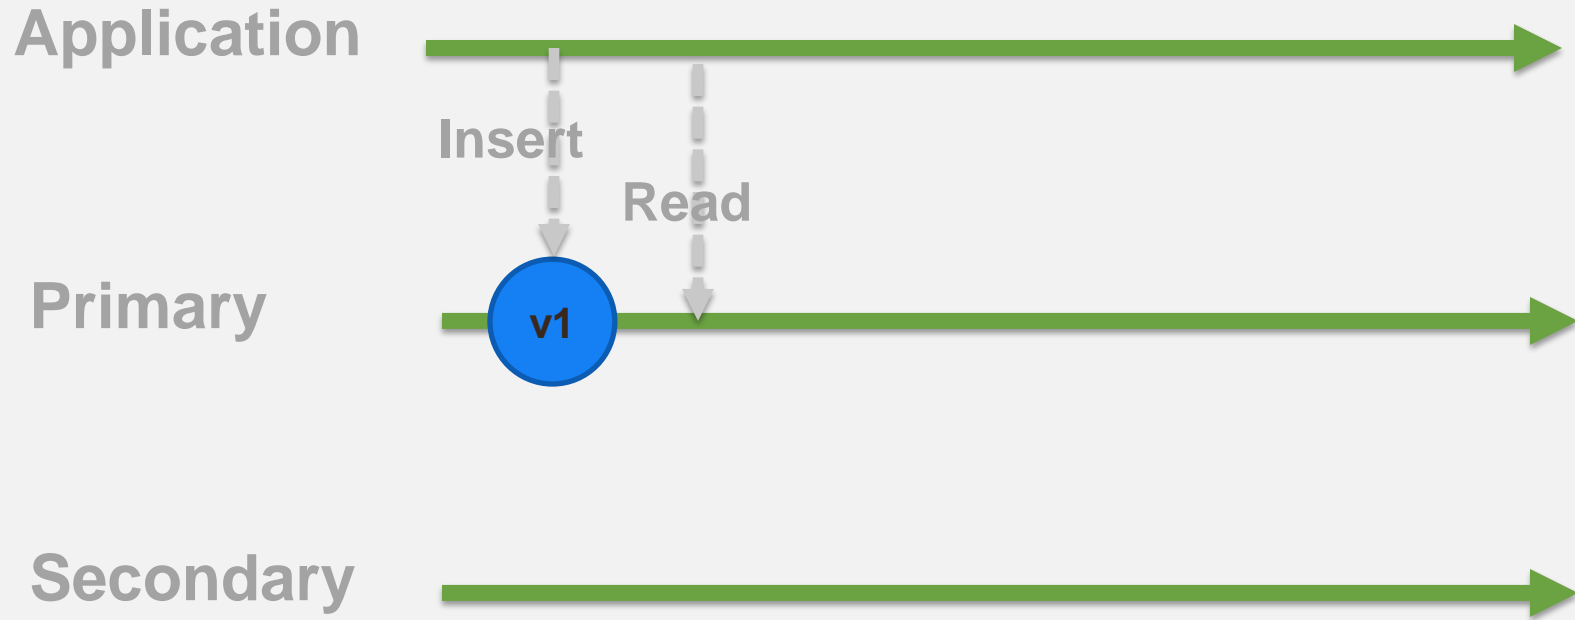
Backup node

```
> rs.initiate(conf)
```

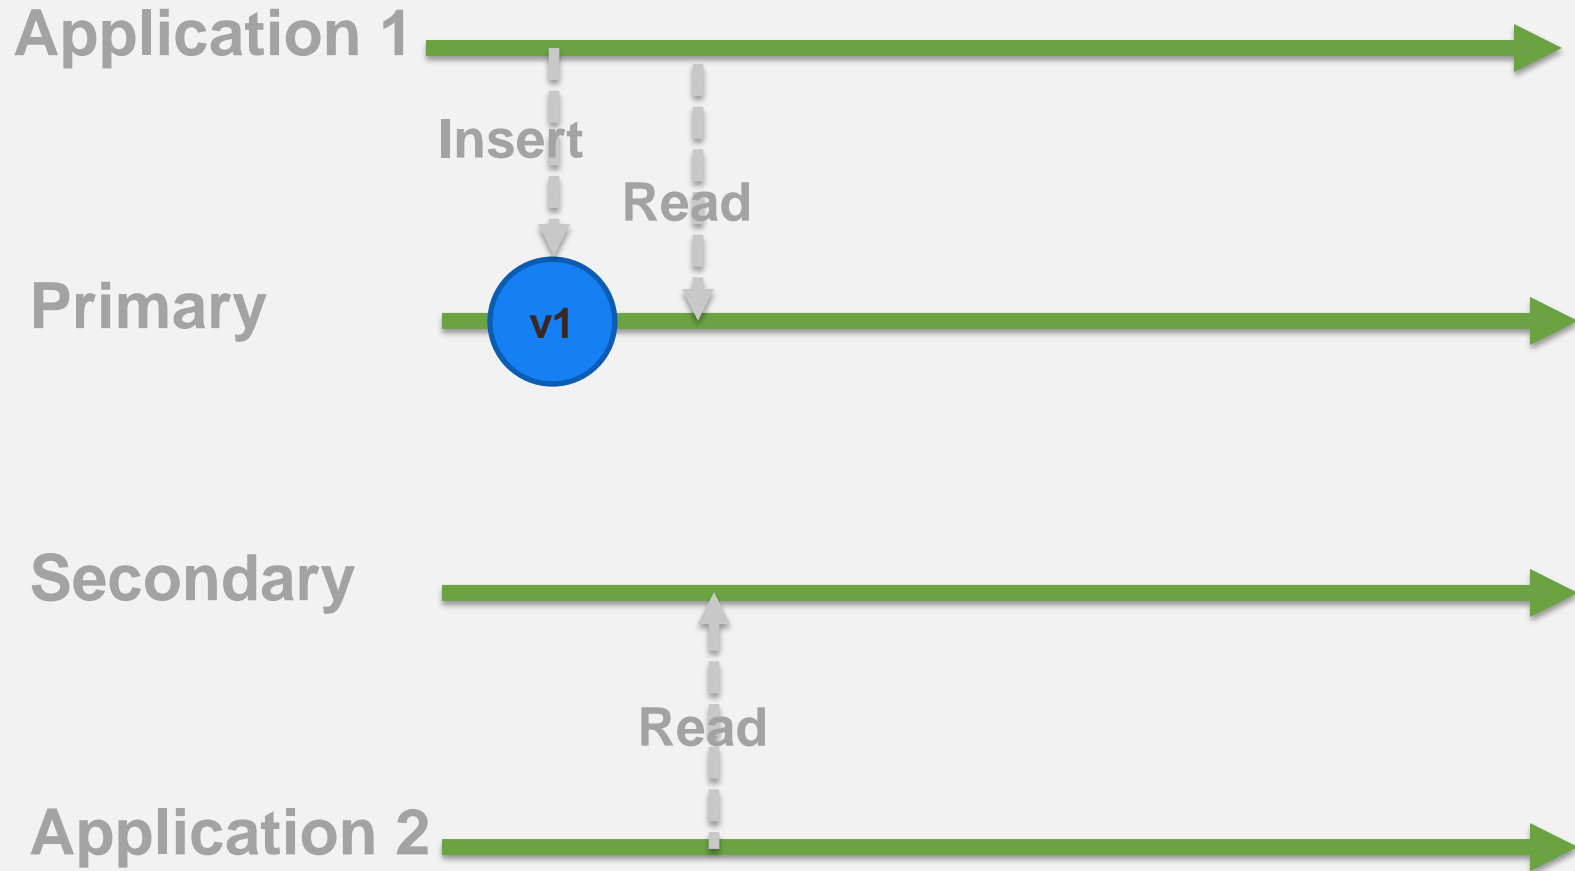
Developing with Replica Sets

- Consistency
- Write Preference
- Read Preference

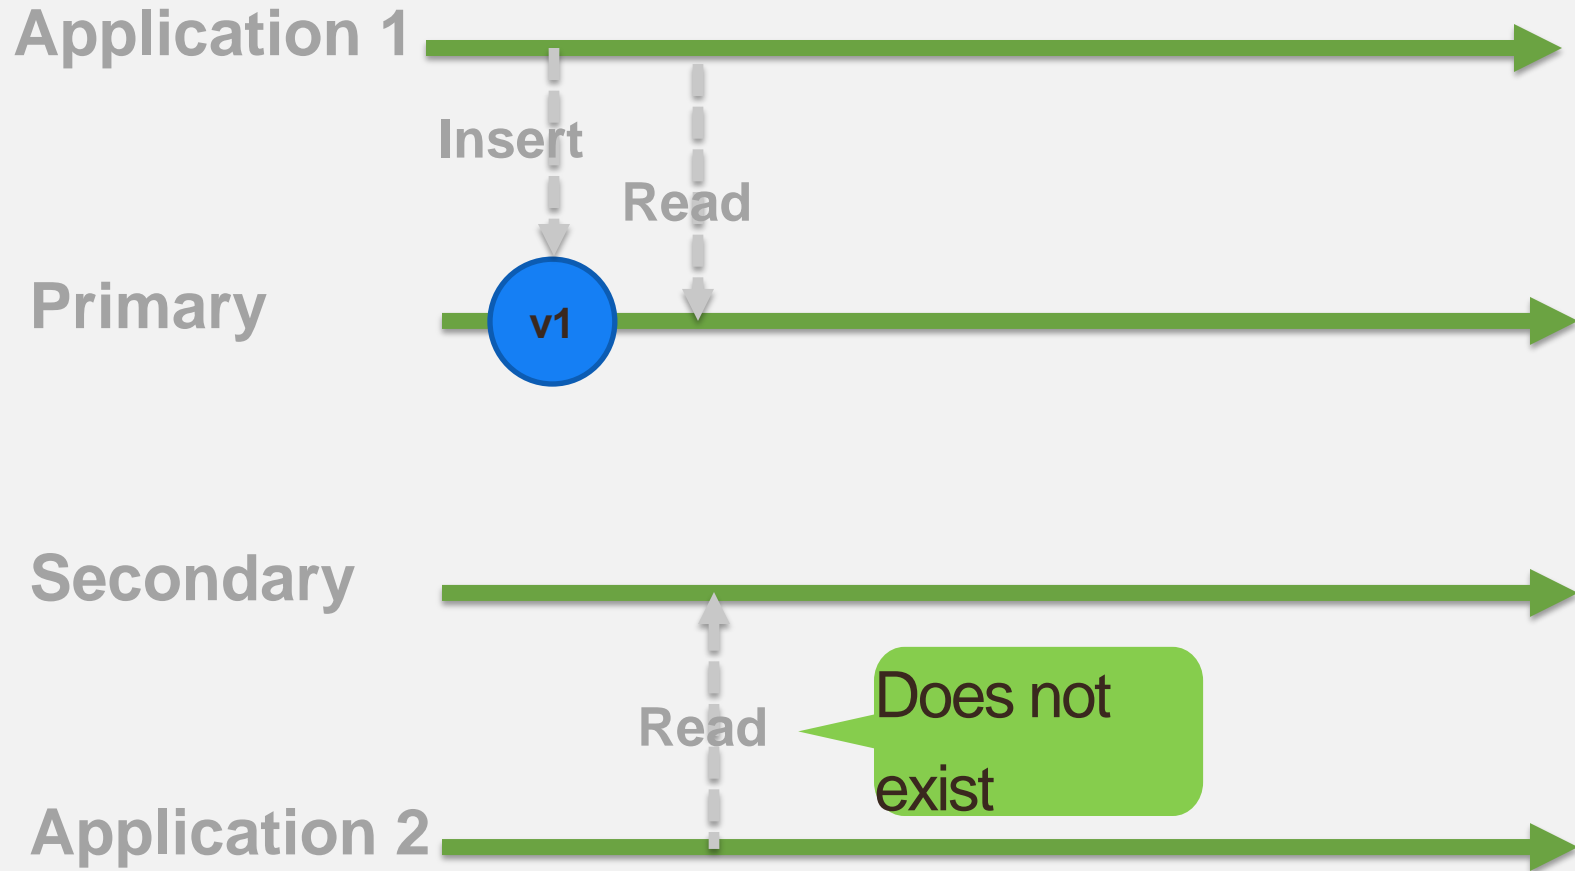
Consistency – Strong



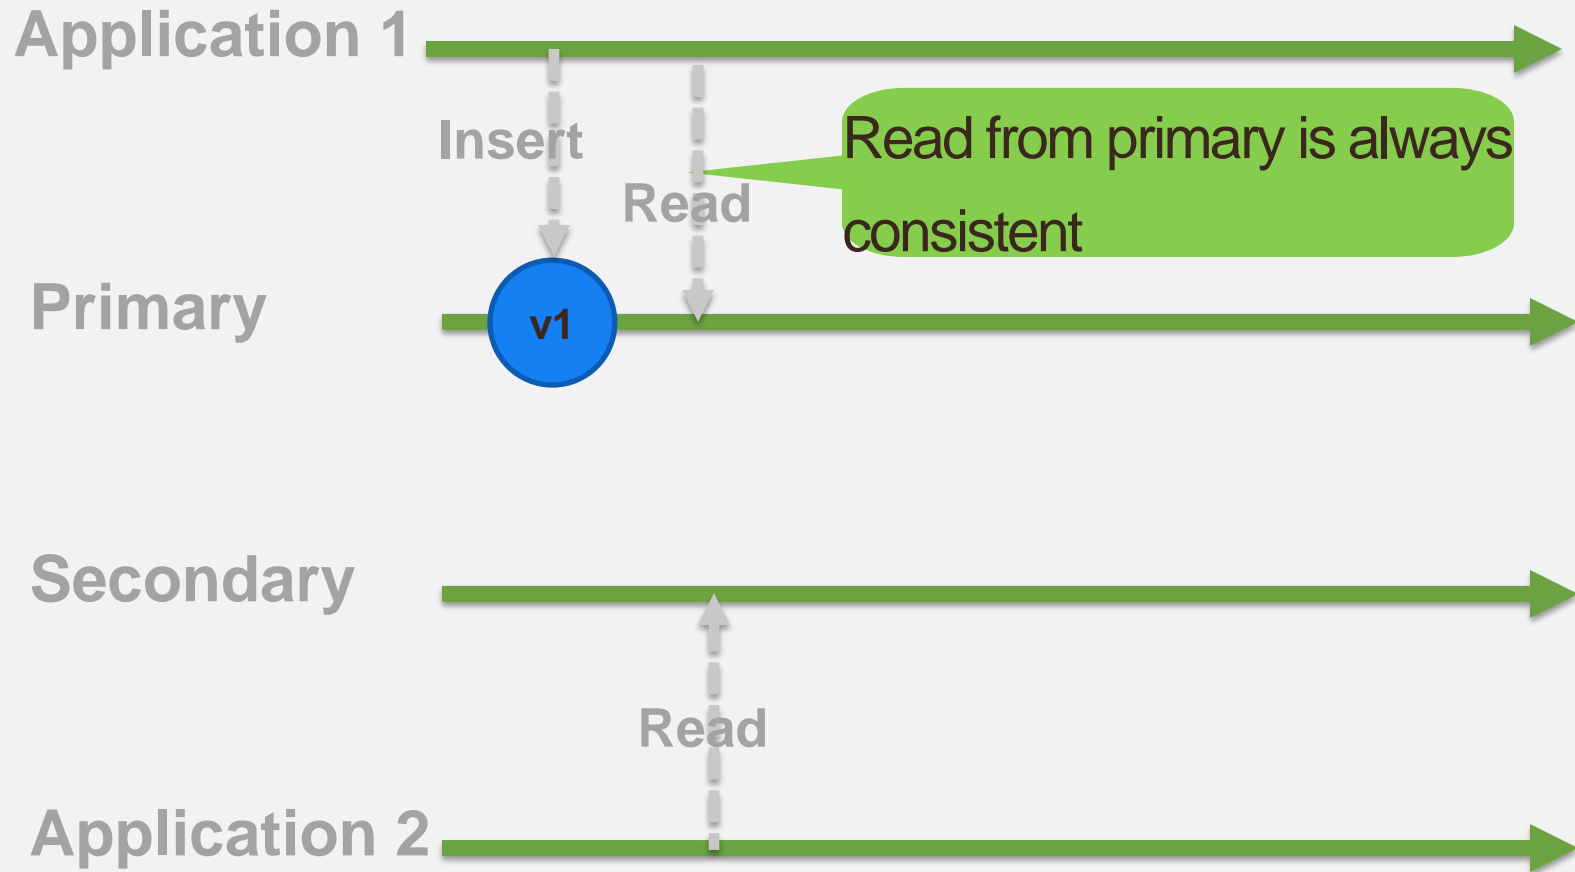
Consistency – Delayed




Consistency – Delayed



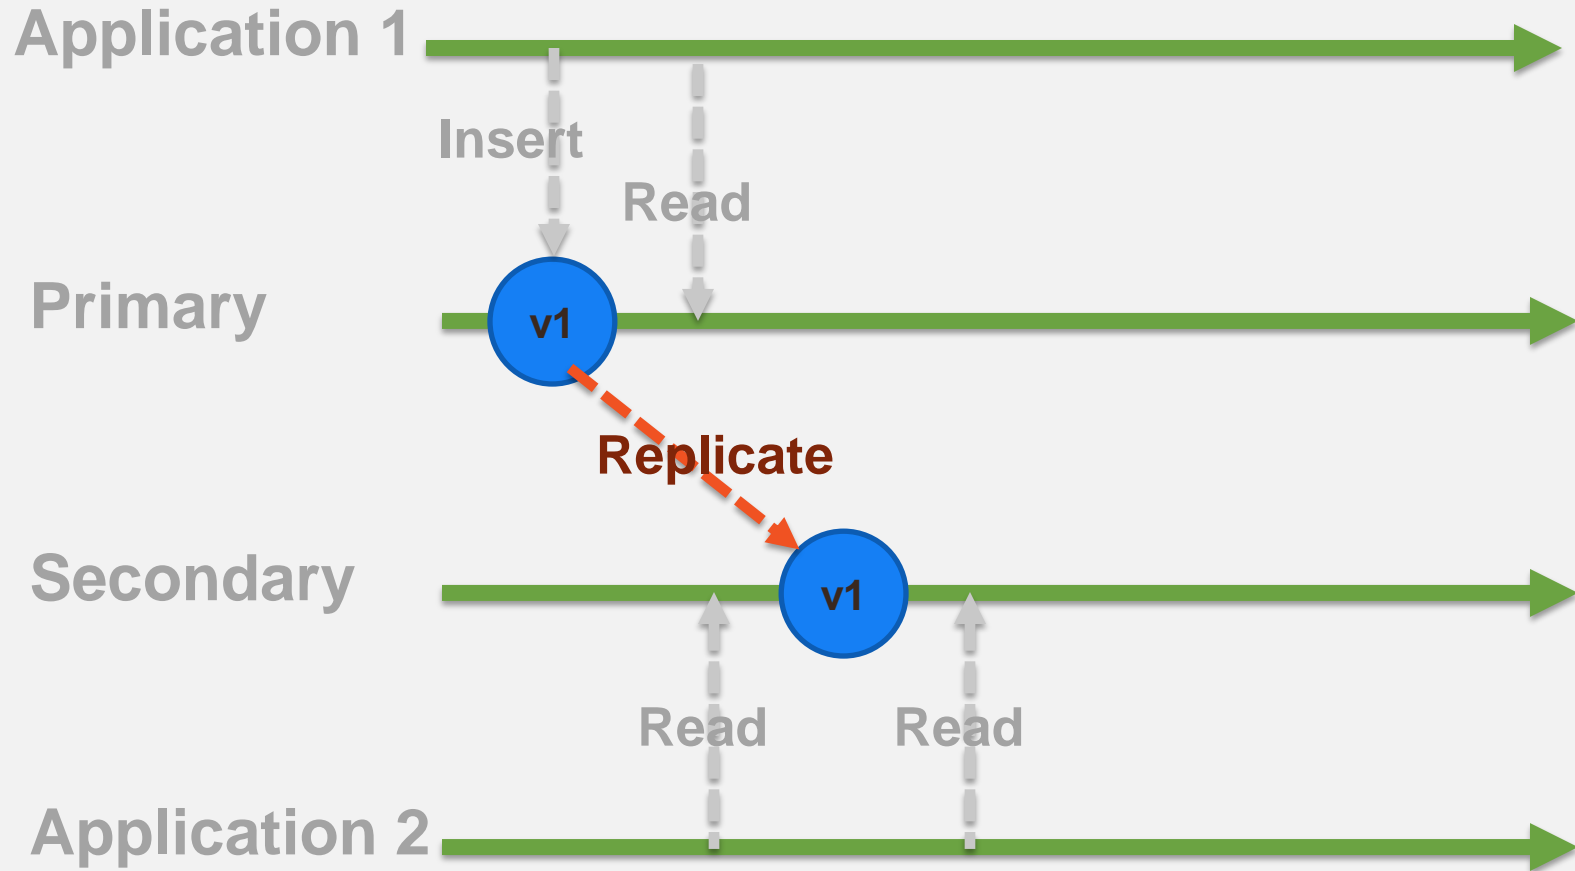
Consistency – Delayed



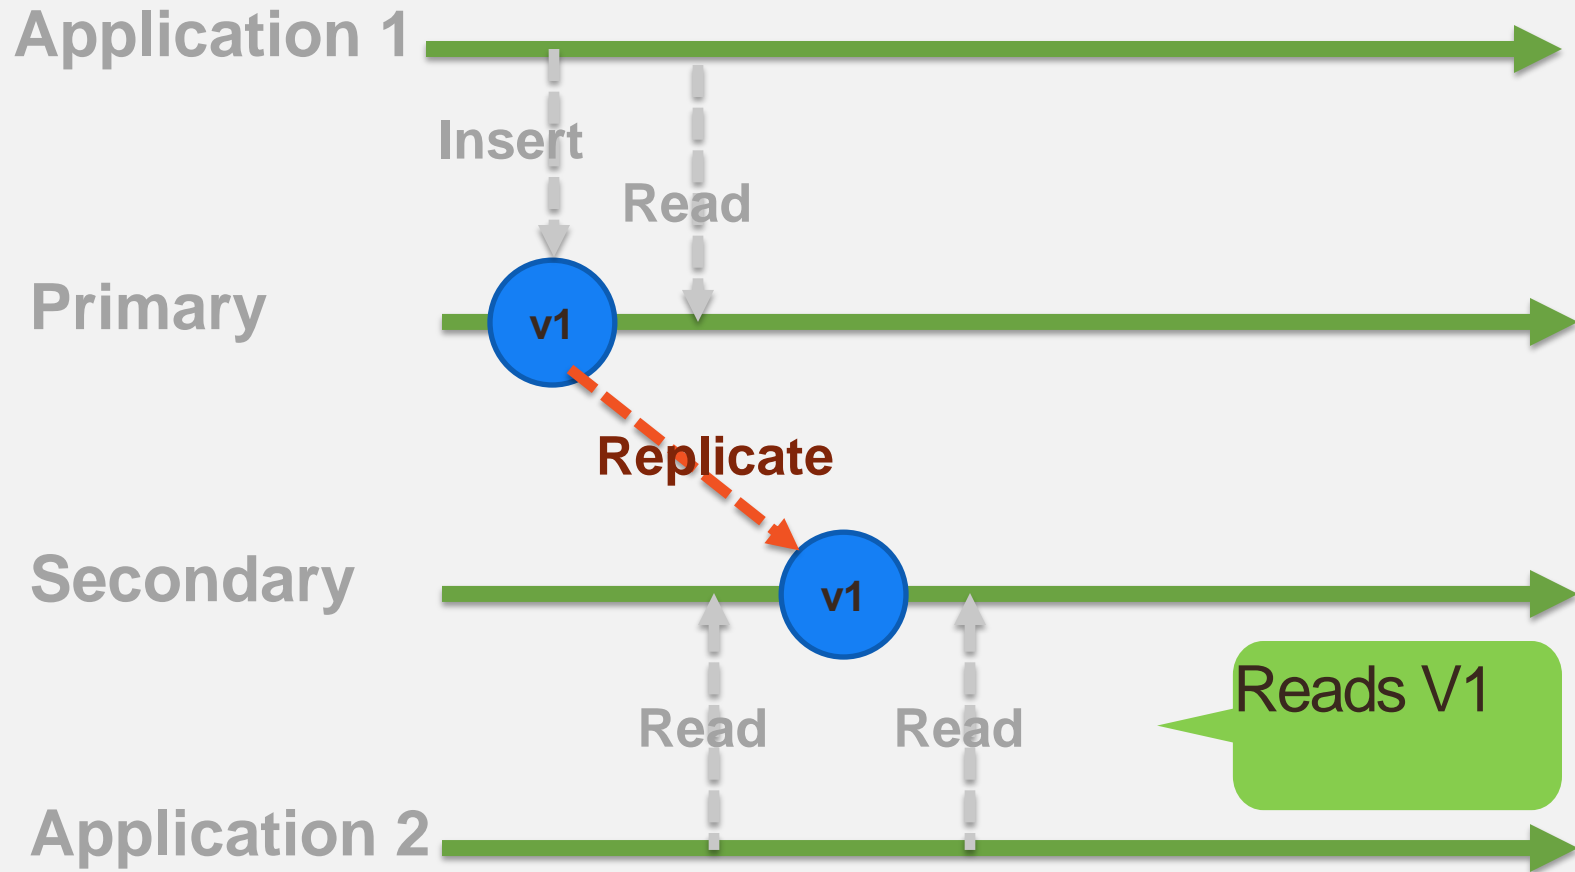
Consistency – Delayed



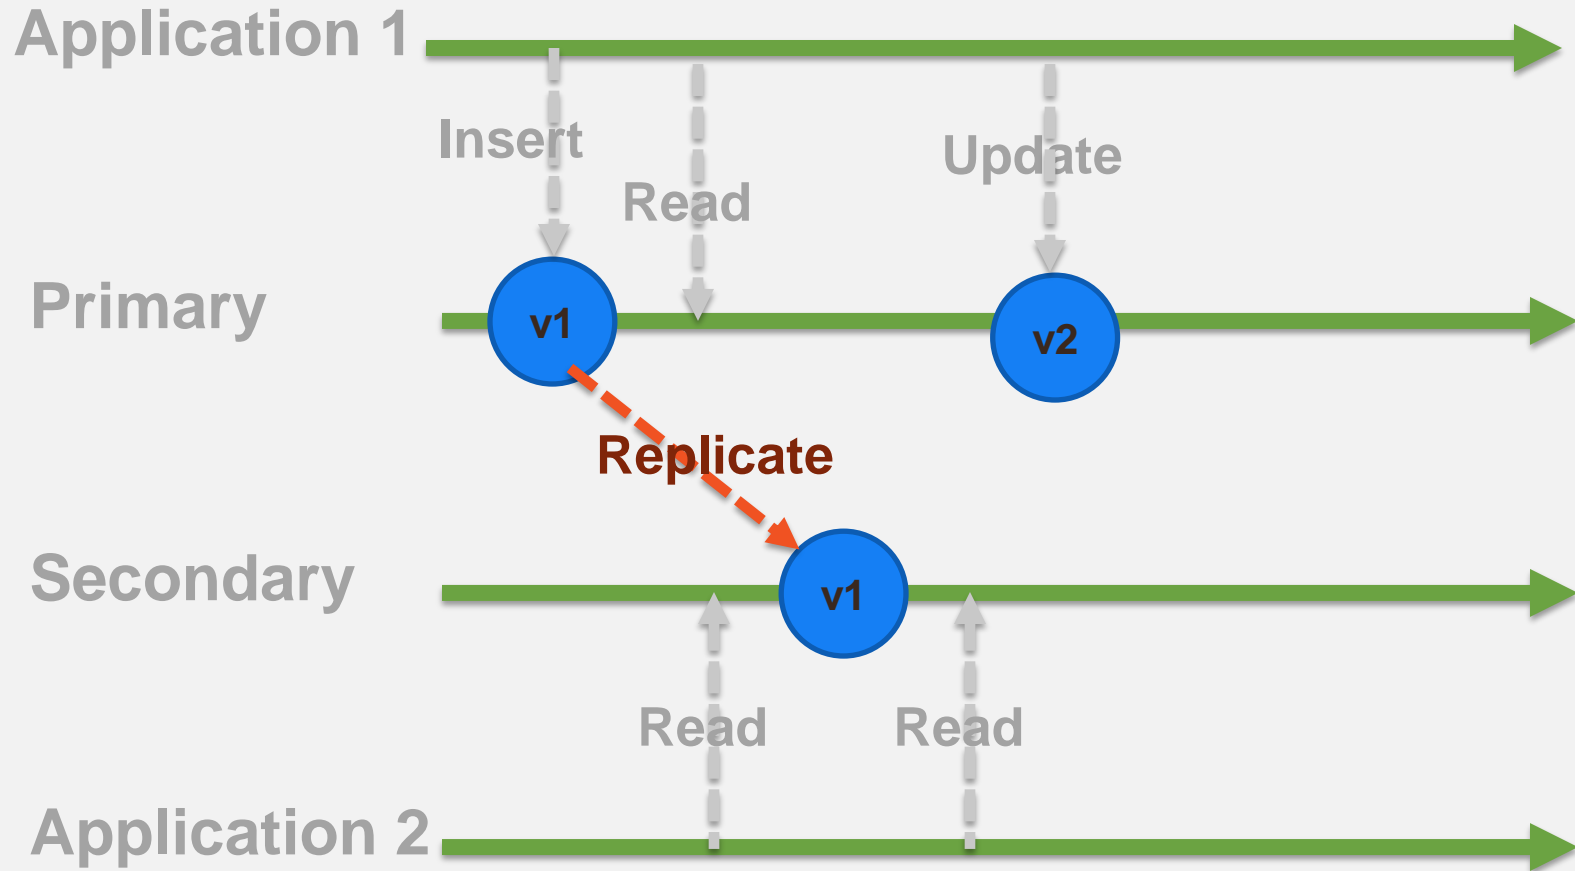
Consistency – Delayed



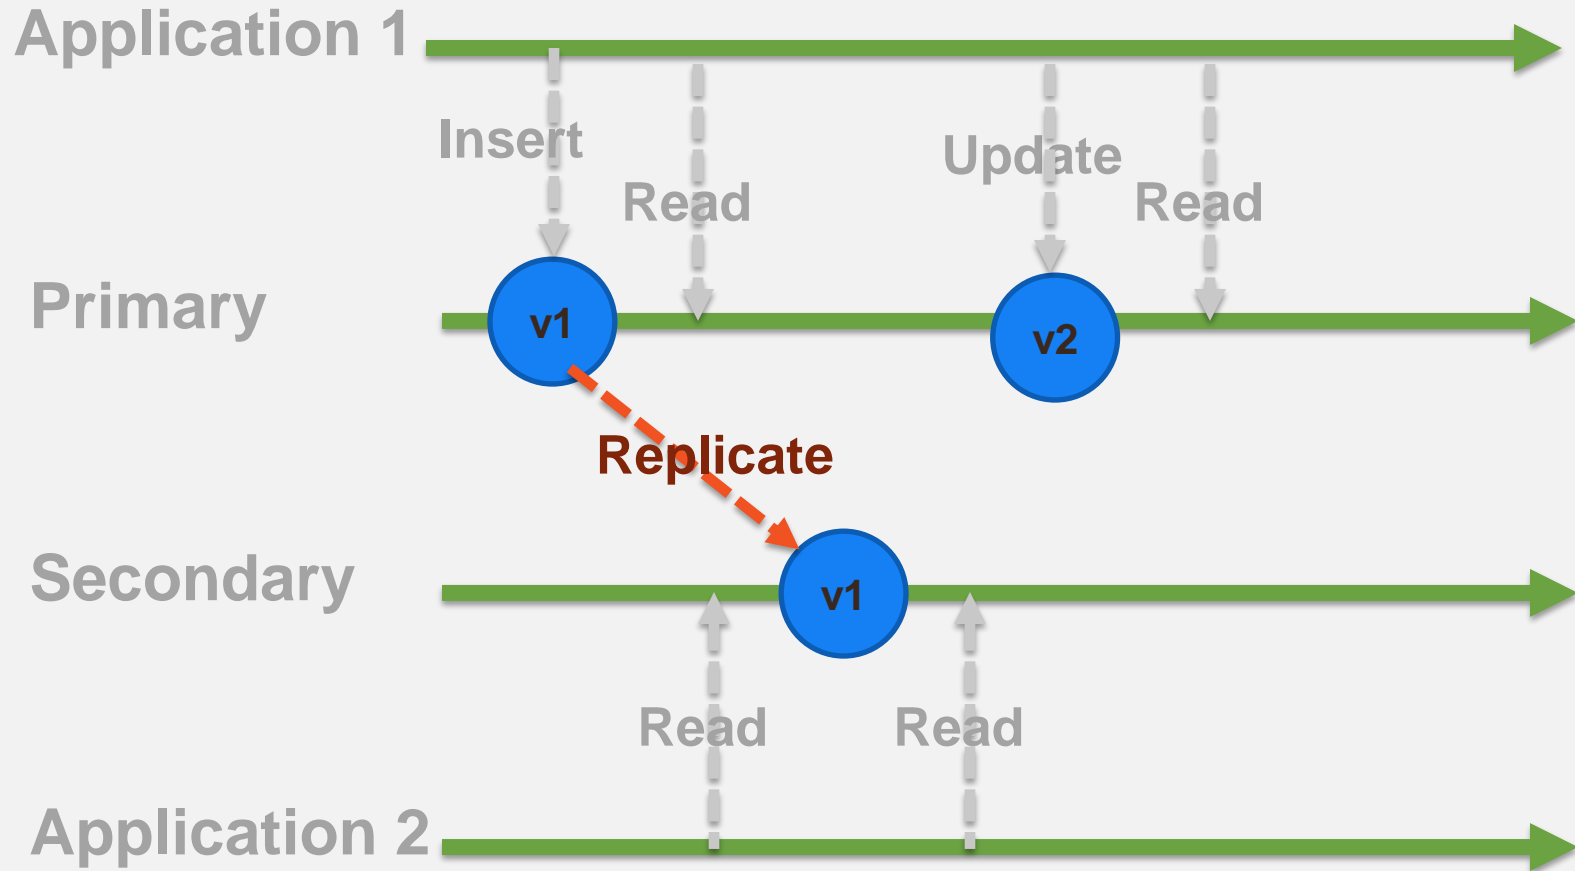
Consistency – Delayed



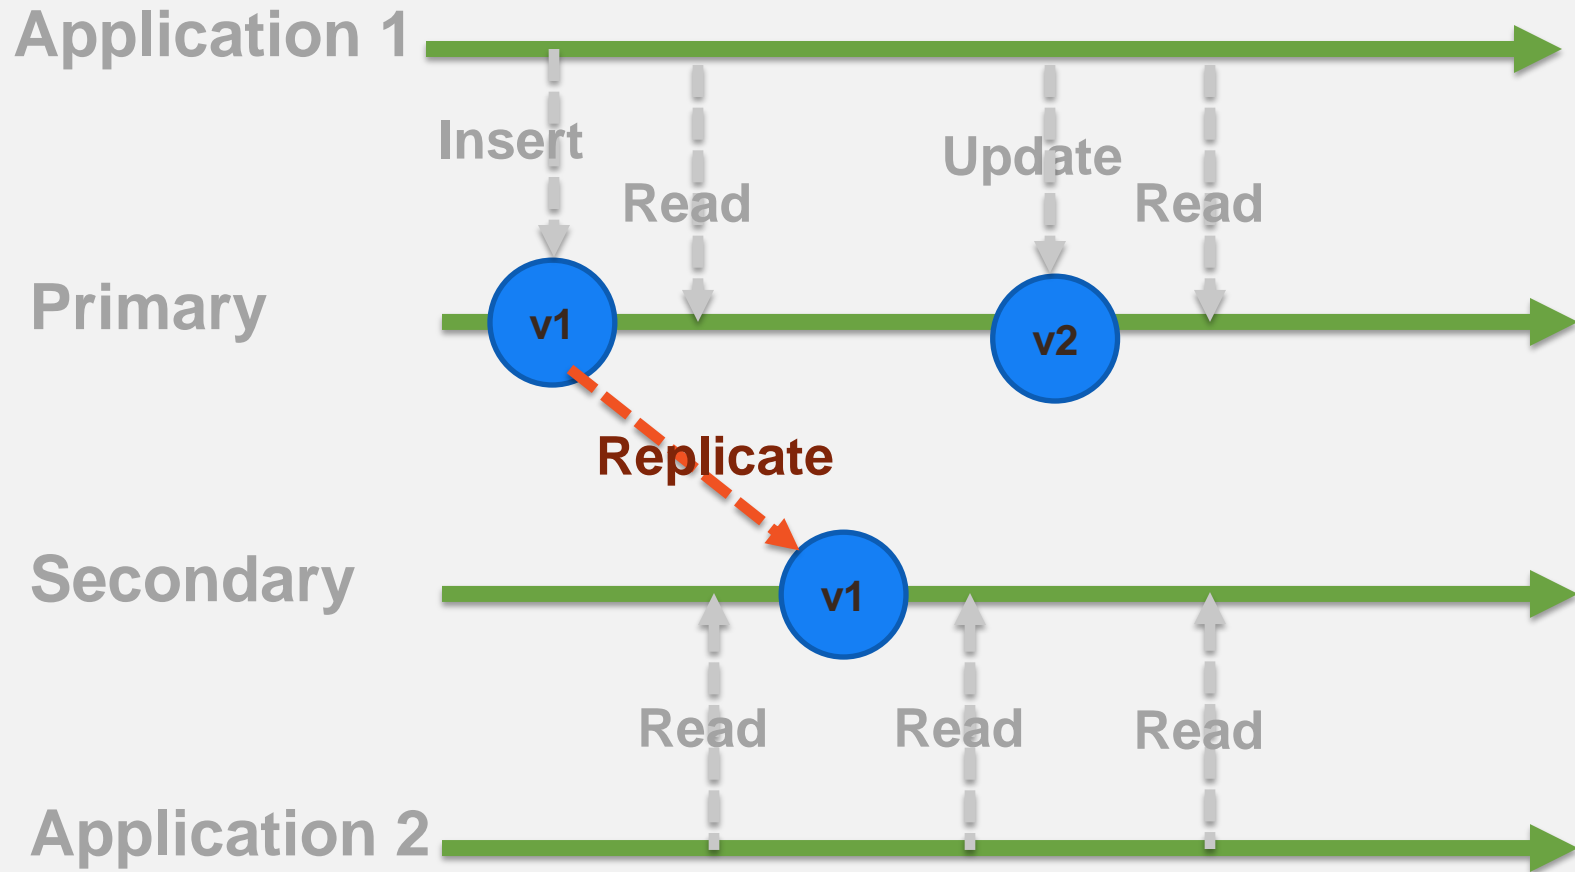
Consistency – Delayed




Consistency – Delayed



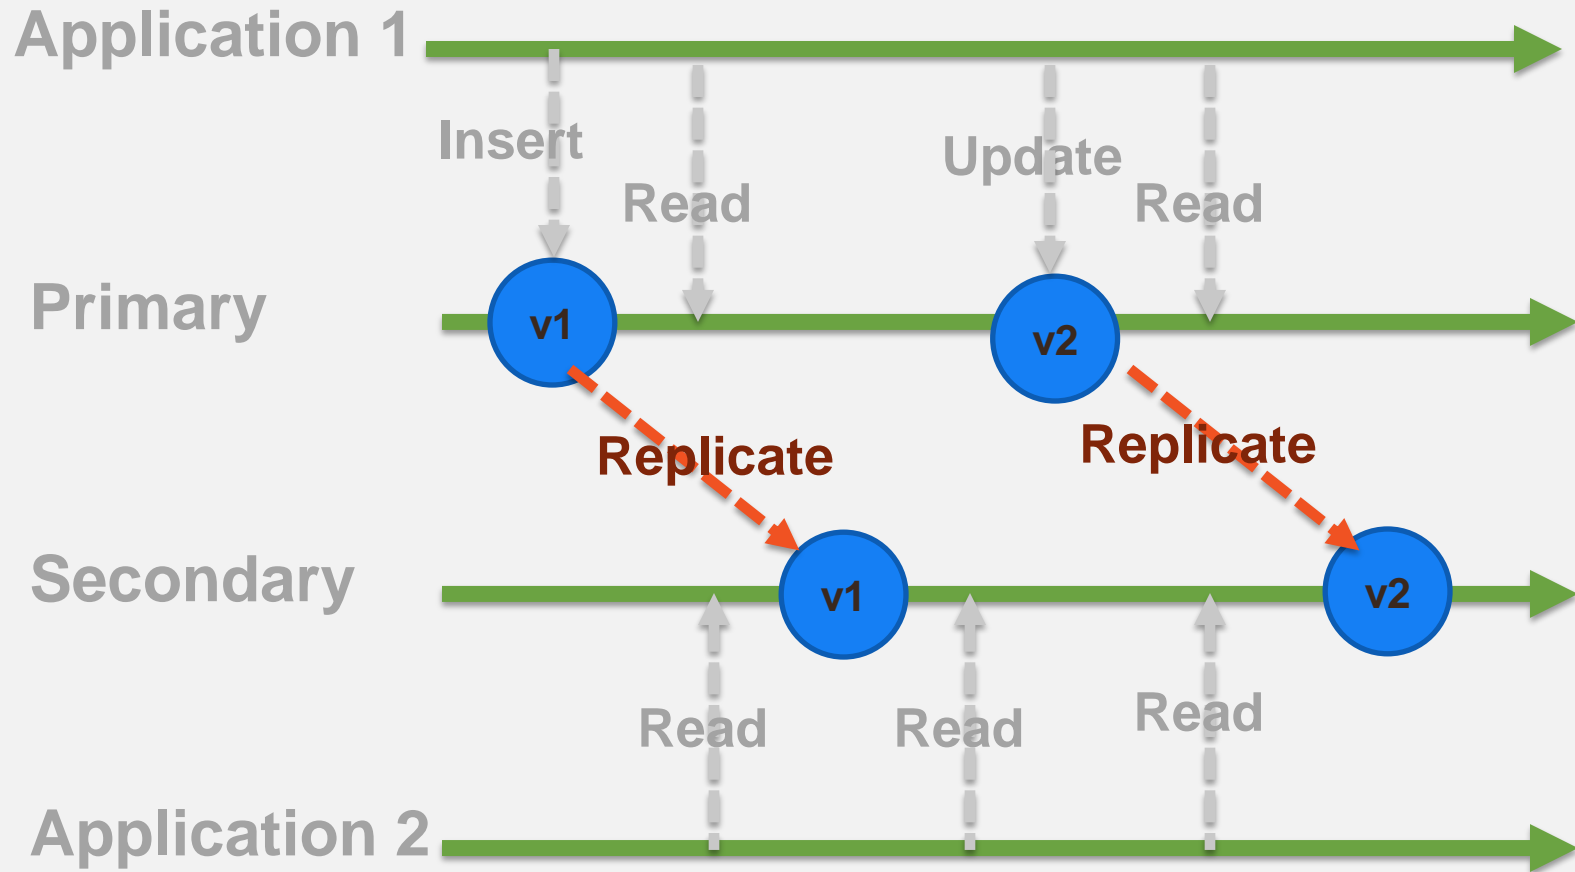
Consistency – Delayed



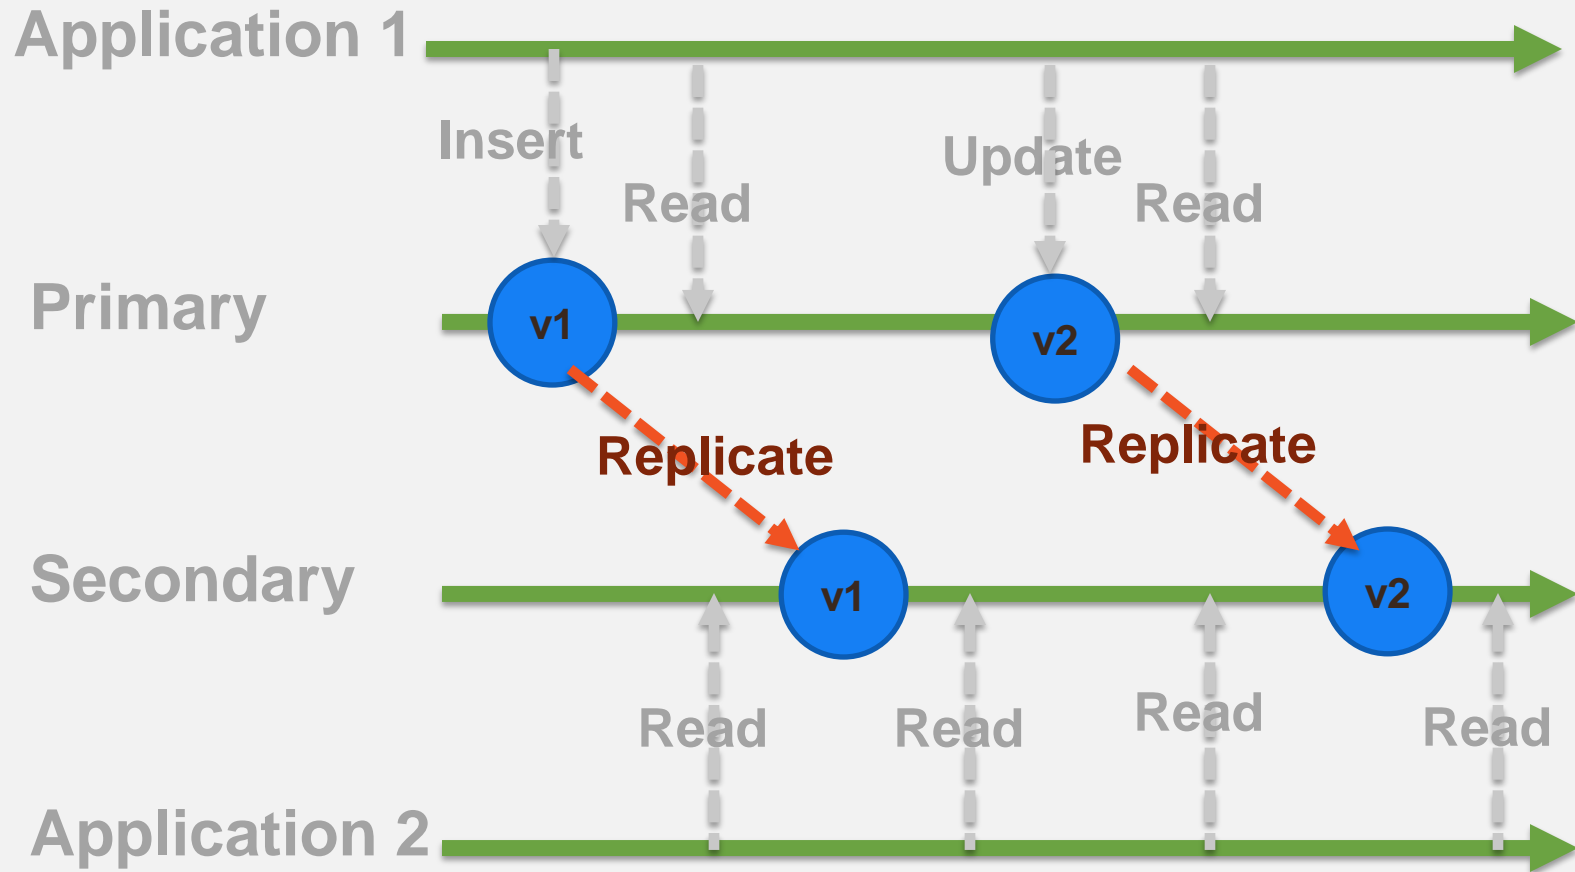
Consistency – Delayed



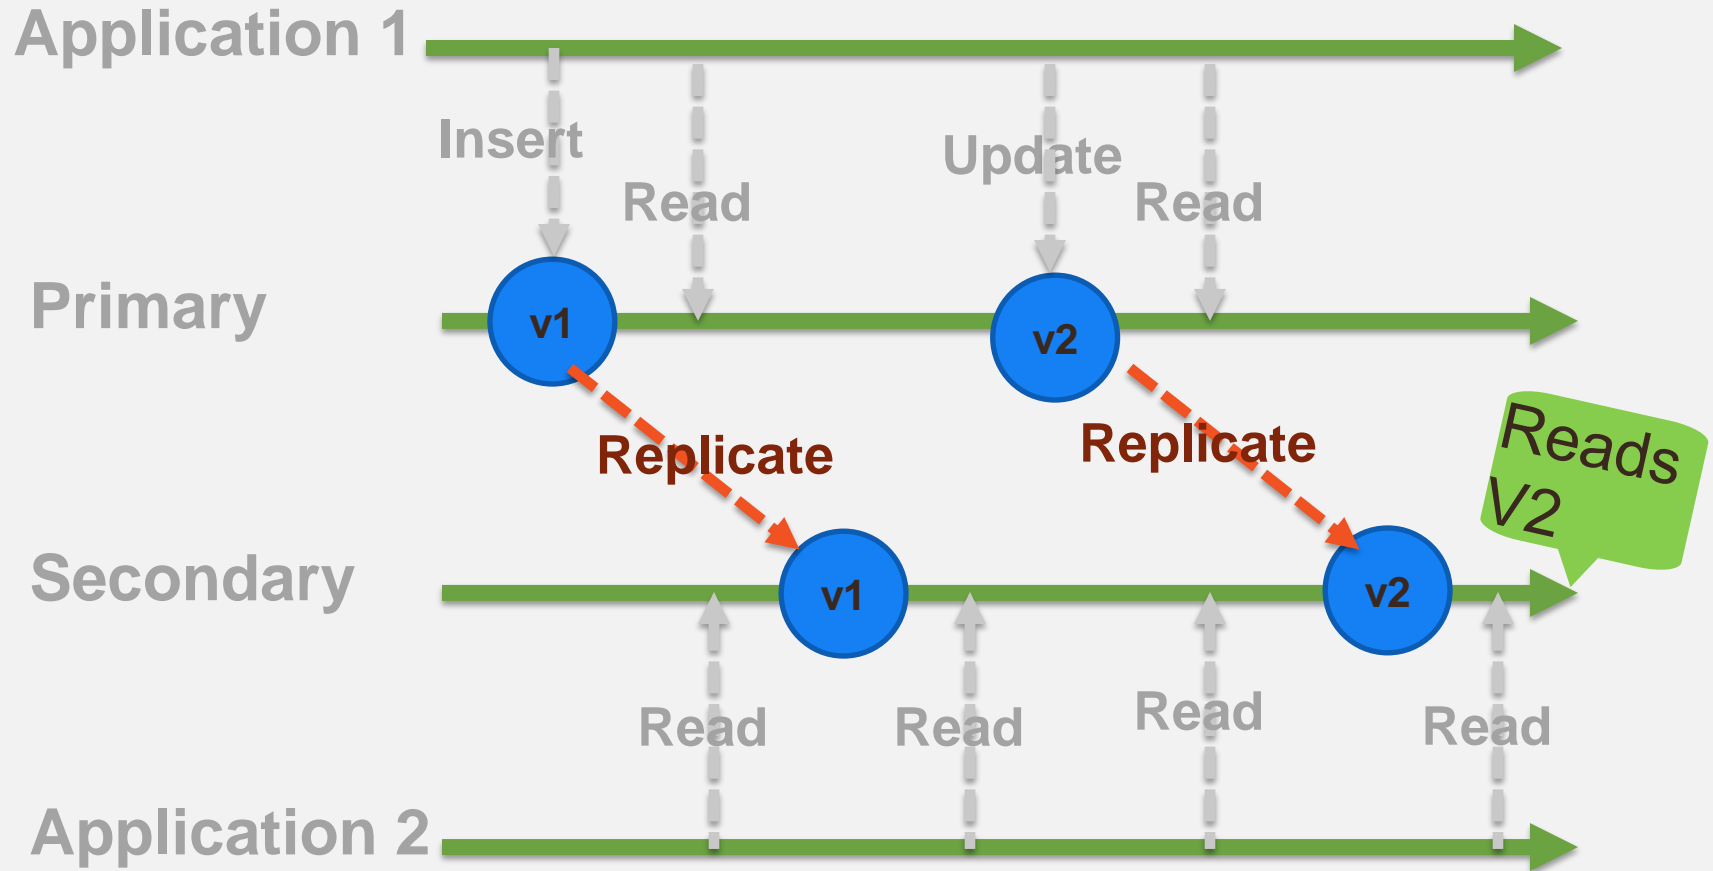
Consistency – Delayed



Consistency – Delayed



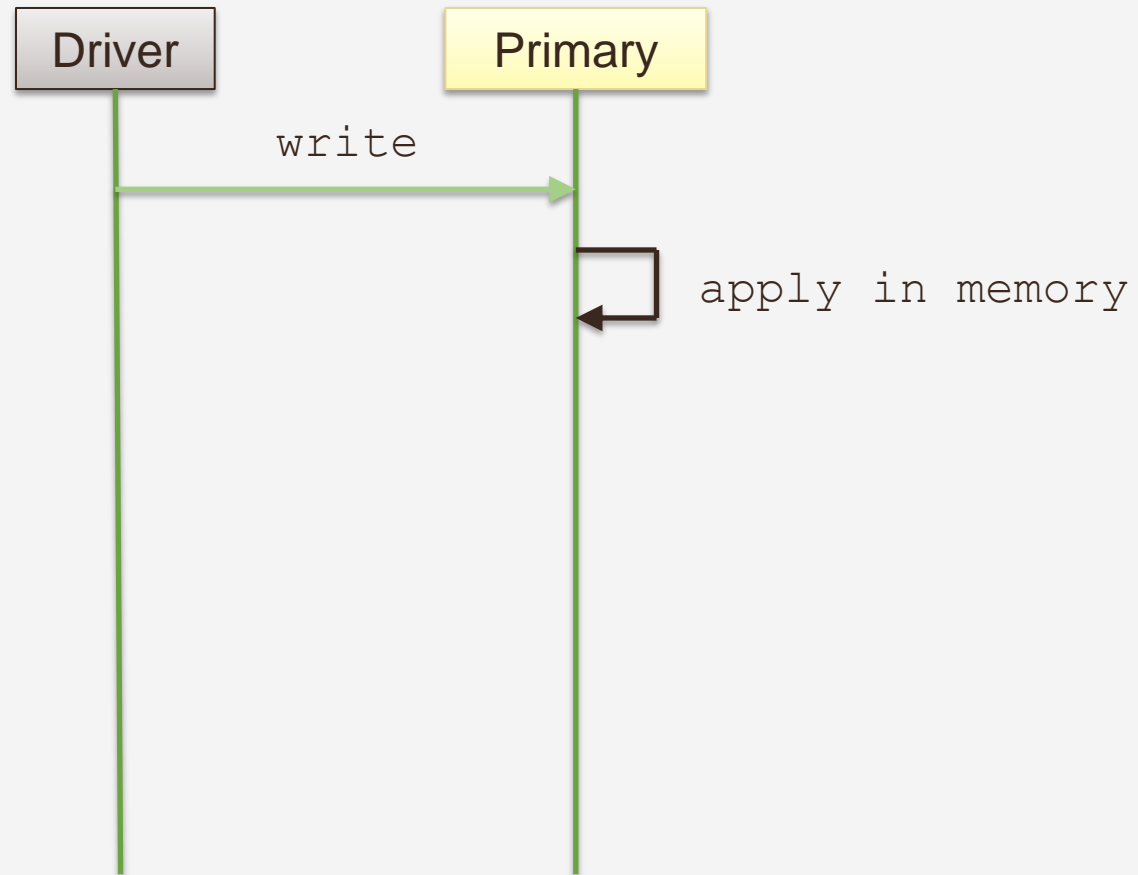
Consistency – Delayed



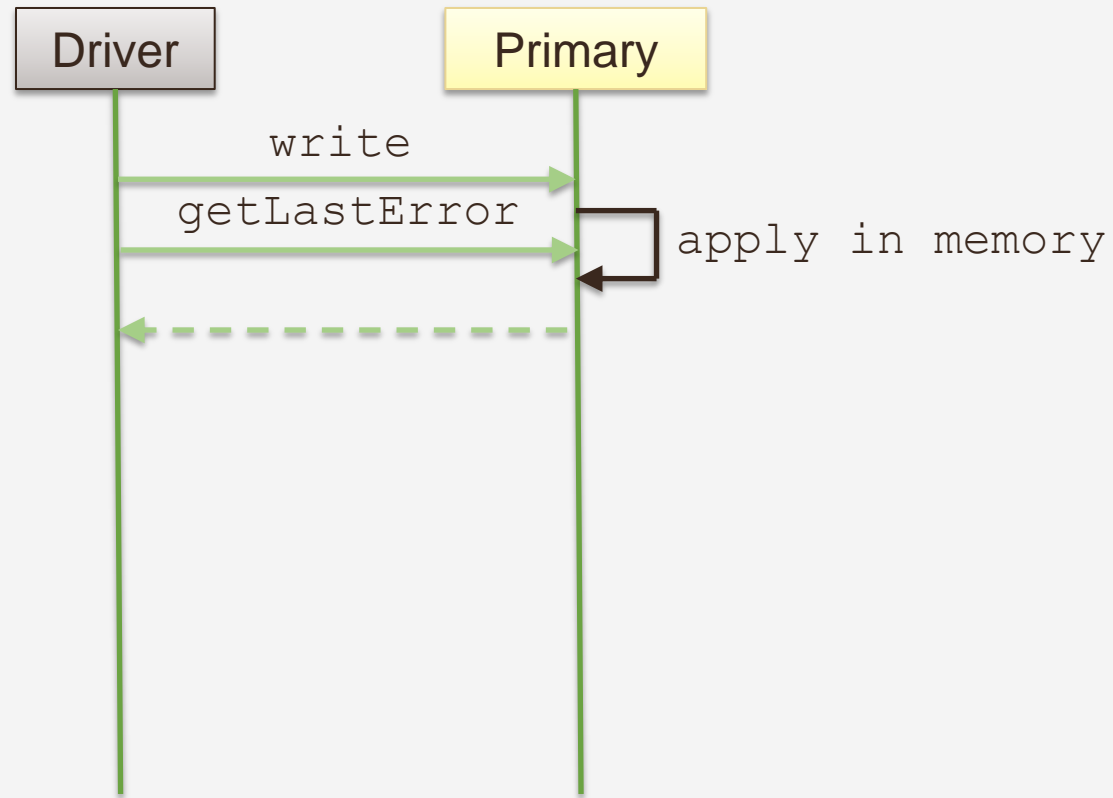
Write Preference

- Network acknowledged
- Wait for error
- Wait for journal sync
- Wait for replication

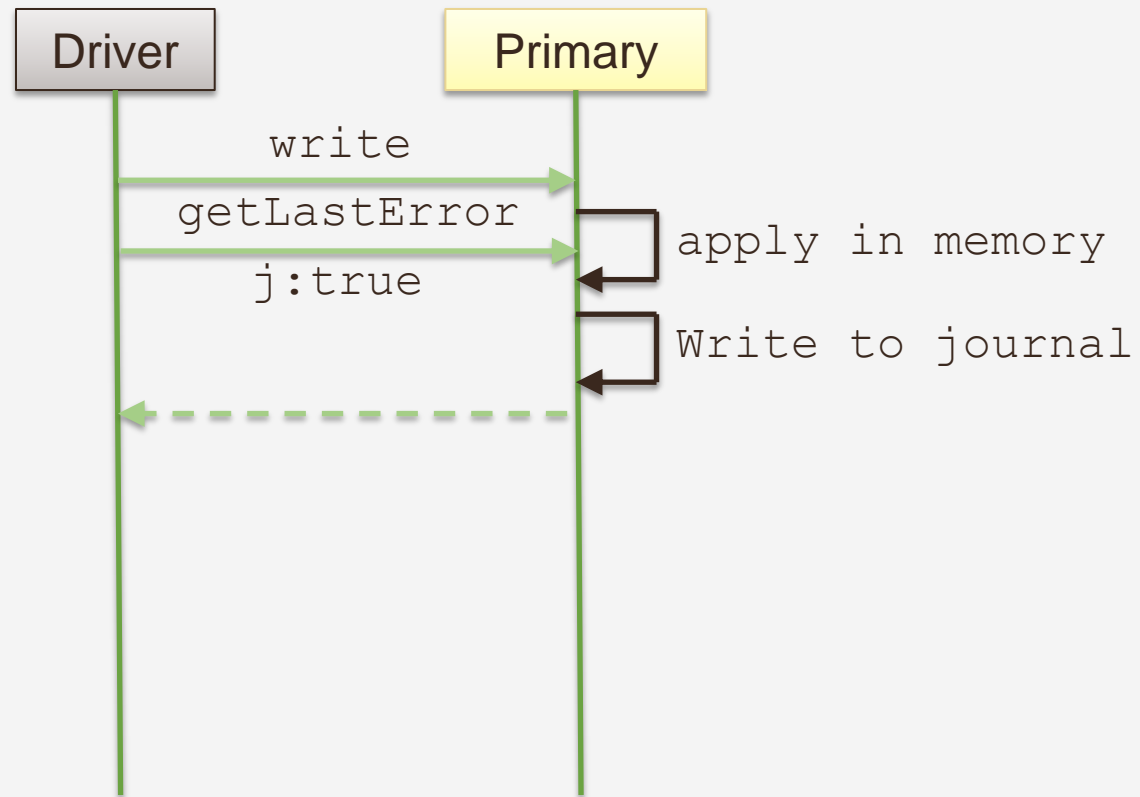
Network Acknowledged



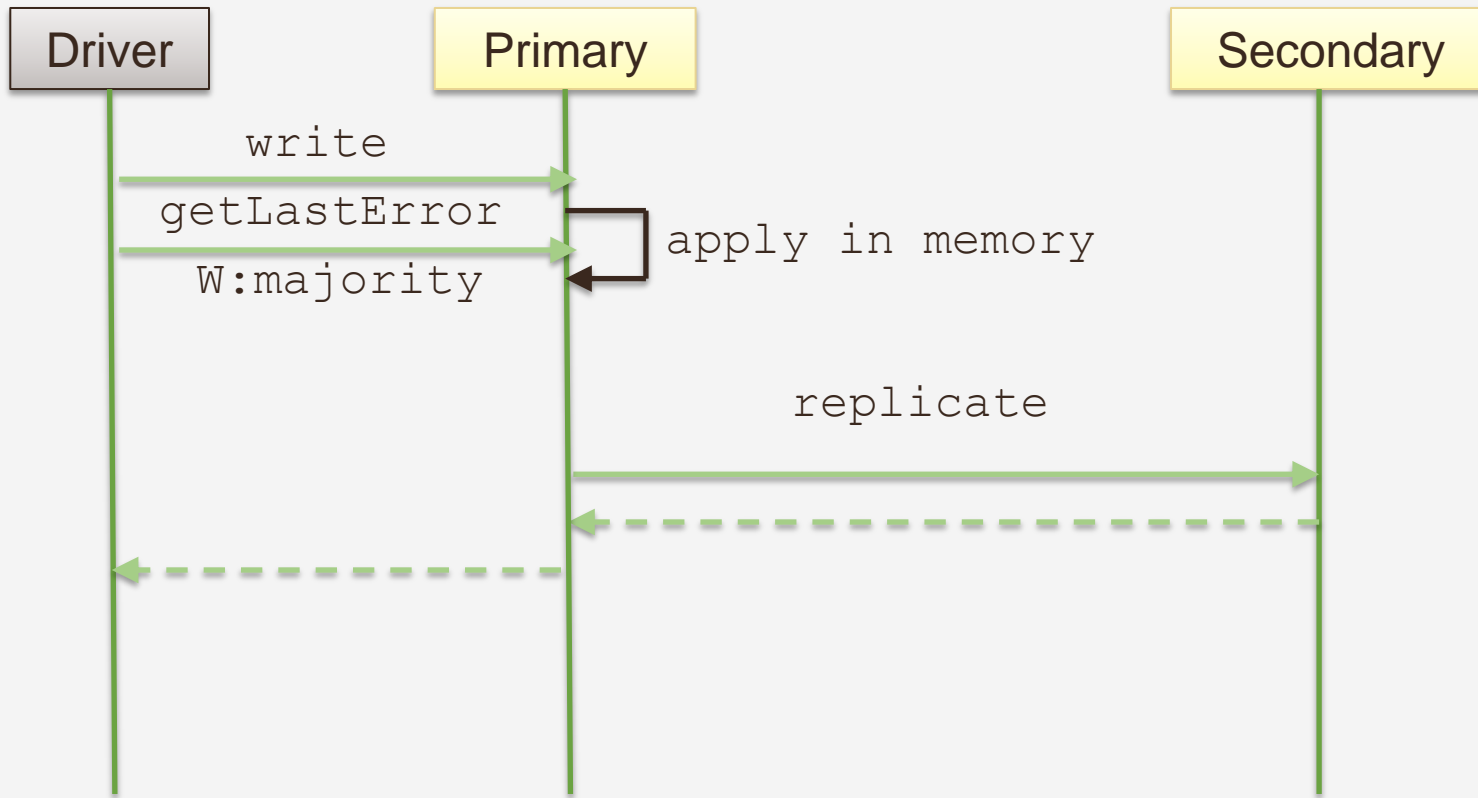
Wait for error



Wait for journal sync



Wait for replication



Tagging

- Since 2.0.0
- Control over where data is written to
- Each member can have one or more tags e.g.
 - Dc : "ny"
 - dc: "ny", ip: "192.168", rack: "row3rk7"
- Replica set defines rules for where data resides
- Rules can change without changing app code

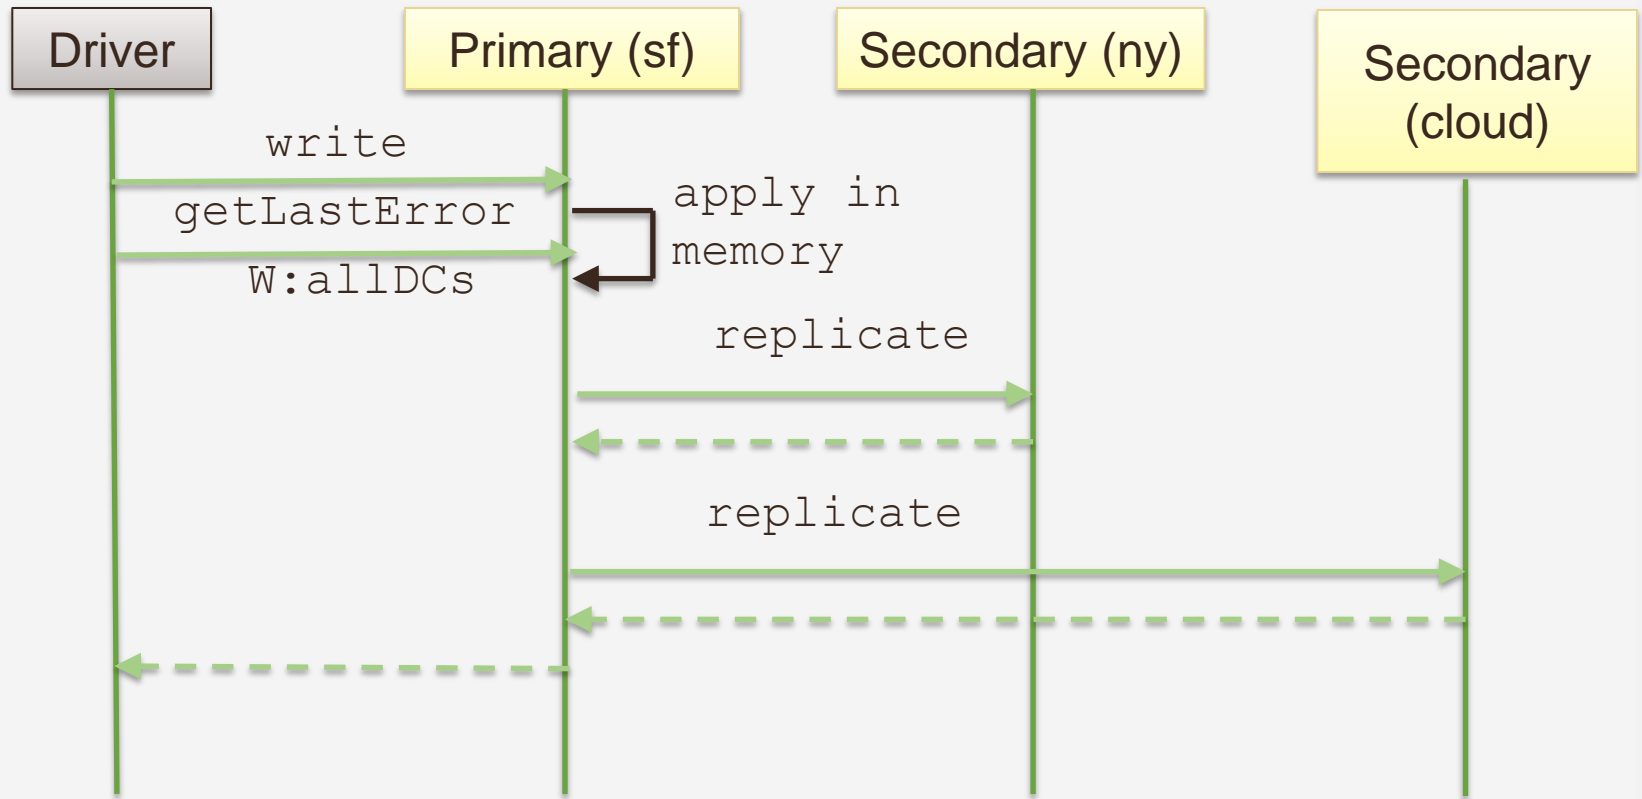
Tagging - example

```
{
  _id : "mySet",
  members : [
    { _id : 0, host : "A", tags : { "dc": "ny" } },
    { _id : 1, host : "B", tags : { "dc": "ny" } },
    { _id : 2, host : "C", tags : { "dc": "sf" } },
    { _id : 3, host : "D", tags : { "dc": "sf" } },
    { _id : 4, host : "E", tags : { "dc": "cloud" } } ]
  settings : {
    getLastErrorModes : {
      allDCs : { "dc" : 3 },
      someDCs : { "dc" : 2 } }
  }
}
```

```
> db.blogs.insert({...})
```

```
> db.runCommand({getError : 1, w : "allDCs"})
```


Wait for replication with tags



Read Preference

- 5 modes (new in 2.2)
 - PRIMARY(only) - Default
 - PRIMARYPREFERRED
 - SECONDARY (only)
 - SECONDARYPREFERRED
 - NEAREST

Tag sets

- Custom read preferences
- Control where you read from
 - E.g. { "disk": "ssd", "use": "reporting" }
- Use in conjunction with standard read preferences
 - Except primary

Operational Considerations

- Upgrade/Maintenance
- Common Deployment Scenarios

Maintenance and Upgrade

- No downtime
- Rolling upgrade/maintenance
 - Start with Secondary
 - Primary last

Replica Set – 1 Data Center




- Single datacenter
- Single switch & power
- Points of failure:
 - Power
 - Network
 - Datacenter
 - Two node failure
- Automatic recovery of single node crash

Replica Set – 2 data centers



- Multi datacenter
- DR node for safety
- Can't do multi data center durable write safely since only 1 node in distant DC

Replica Set – 3 Data Centers



- Three data centers
- Can survive full data center loss
- Can do $w = \{ dc : 2 \}$ to guarantee write in 2 data centers (with tags)

Behind the Curtain

- Schema
- Oplog

Schema

- Local DB (not replicated)
 - system.replset
 - oplog.rs
 - Capped collection
 - Idempotent version of operation stored

Detections

- Heartbeat every 2 seconds
 - Times out in 10 seconds
- Missed heartbeat considered node down

Oplog

```
> db.replsettest.insert({_id:1,value:1})
```

```
{ "ts" : Timestamp(1350539727000, 1), "h" :  
NumberLong("6375186941486301201"), "op" : "i",  
"ns" : "test.replsettest", "o" : { "_id" : 1,  
"value" : 1 } }
```

```
>
```

```
db.replsettest.update({_id:1},{ $inc:{value:10}})
```

```
{ "ts" : Timestamp(1350539786000, 1), "h" :  
NumberLong("5484673652472424968"), "op" : "u",  
"ns" : "test.replsettest", "o2" : { "_id" : 1 },  
"o" : { "$set" : { "value" : 11 } } }
```

Oplog

```
> db.replsettest.update({}, {$set:{name : 'foo'}},
false, true)

{ "ts" : Timestamp(1350540395000, 1), "h" :
NumberLong("-4727576249368135876"), "op" : "u",
"ns" : "test.replsettest", "o2" : { "_id" : 2 },
"o" : { "$set" : { "name" : "foo" } } }

{ "ts" : Timestamp(1350540395000, 2), "h" :
NumberLong("-7292949613259260138"), "op" : "u",
"ns" : "test.replsettest", "o2" : { "_id" : 3 },
"o" : { "$set" : { "name" : "foo" } } }

{ "ts" : Timestamp(1350540395000, 3), "h" :
NumberLong("-1888768148831990635"), "op" : "u",
"ns" : "test.replsettest", "o2" : { "_id" : 1 },
"o" : { "$set" : { "name" : "foo" } } }
```



#codemash

Coming Next: Part 3.2 : Sharding

Sridhar Nanjundeswaran

Engineer, MongoDB Inc.

@snanjund