



#codemash

MongoDB – Zero to Sharding

Part 2.1 : Schema Design

Sridhar Nanjundeswaran

Engineer, MongoDB Inc.

@snanjund

Agenda

- Working with documents
- Schema design by example
- Other common Patterns

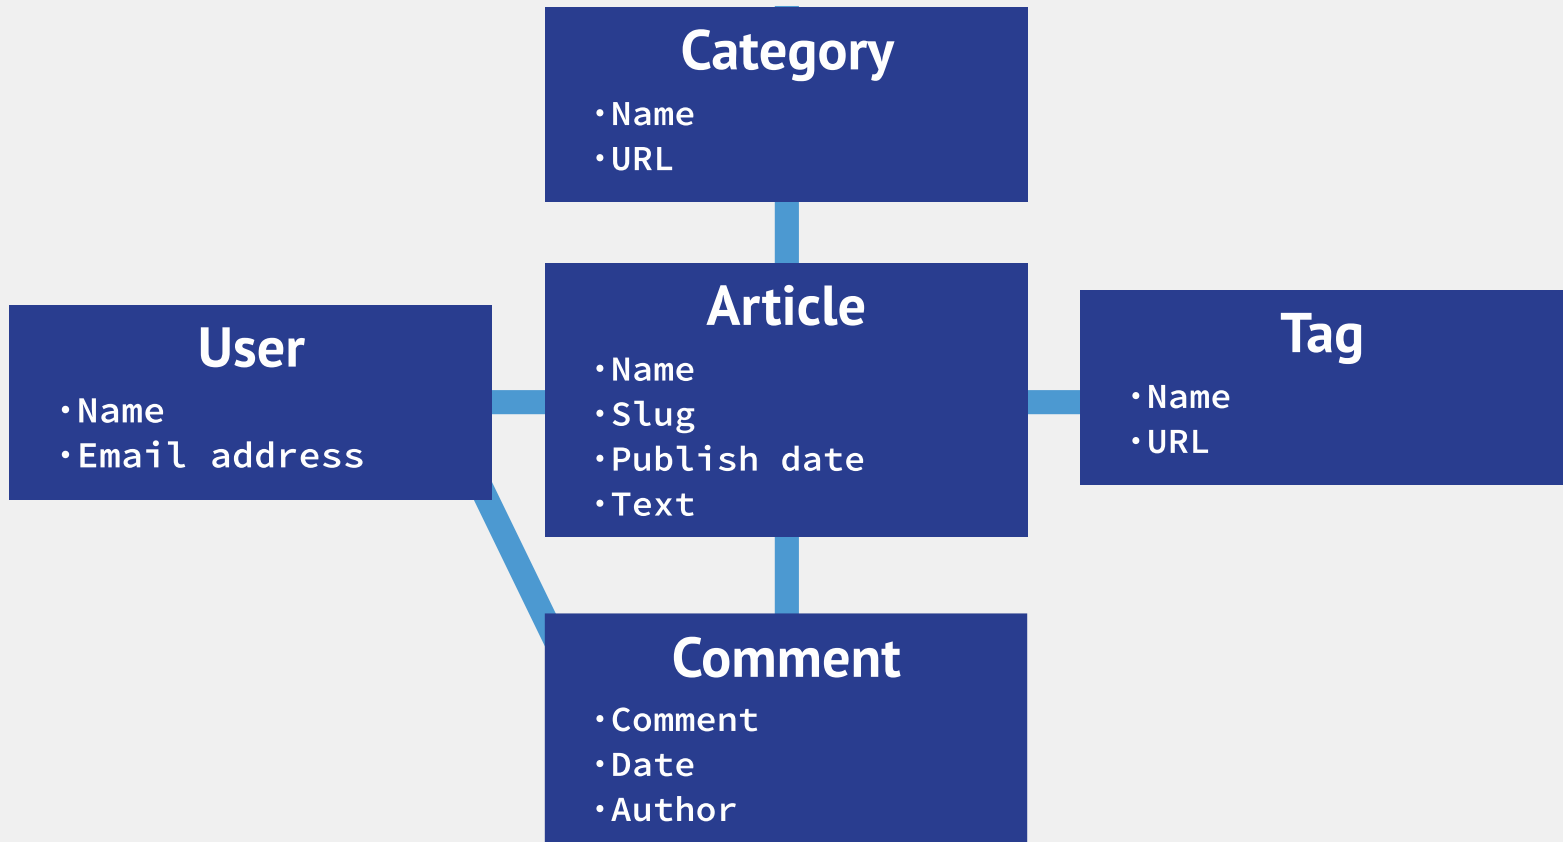
RDBMS		MongoDB
Database	→	Database
Table/View	→	Collection
Row	→	Document
Index	→	Index
Join	→	Embedding & Linking

Terminology

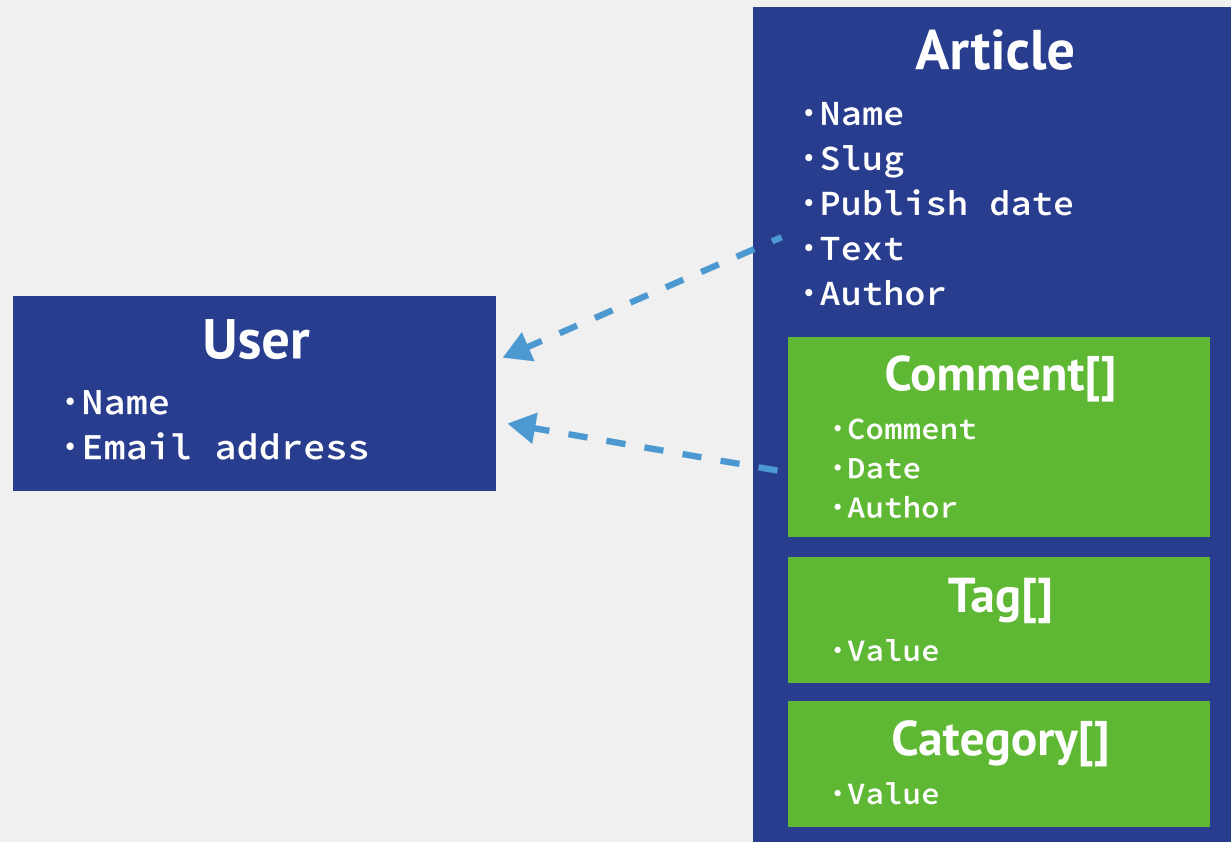
Working with Documents

Documents

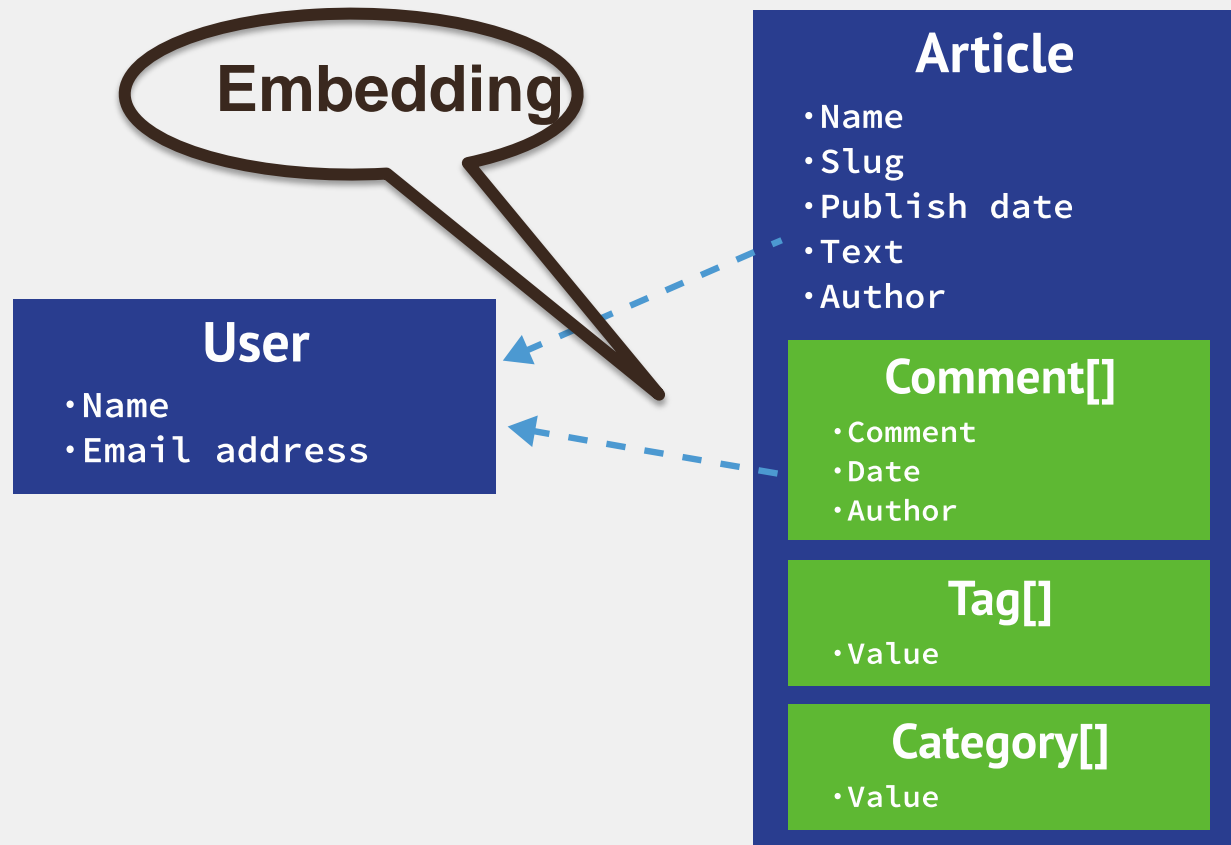
**Provide flexibility and
performance**



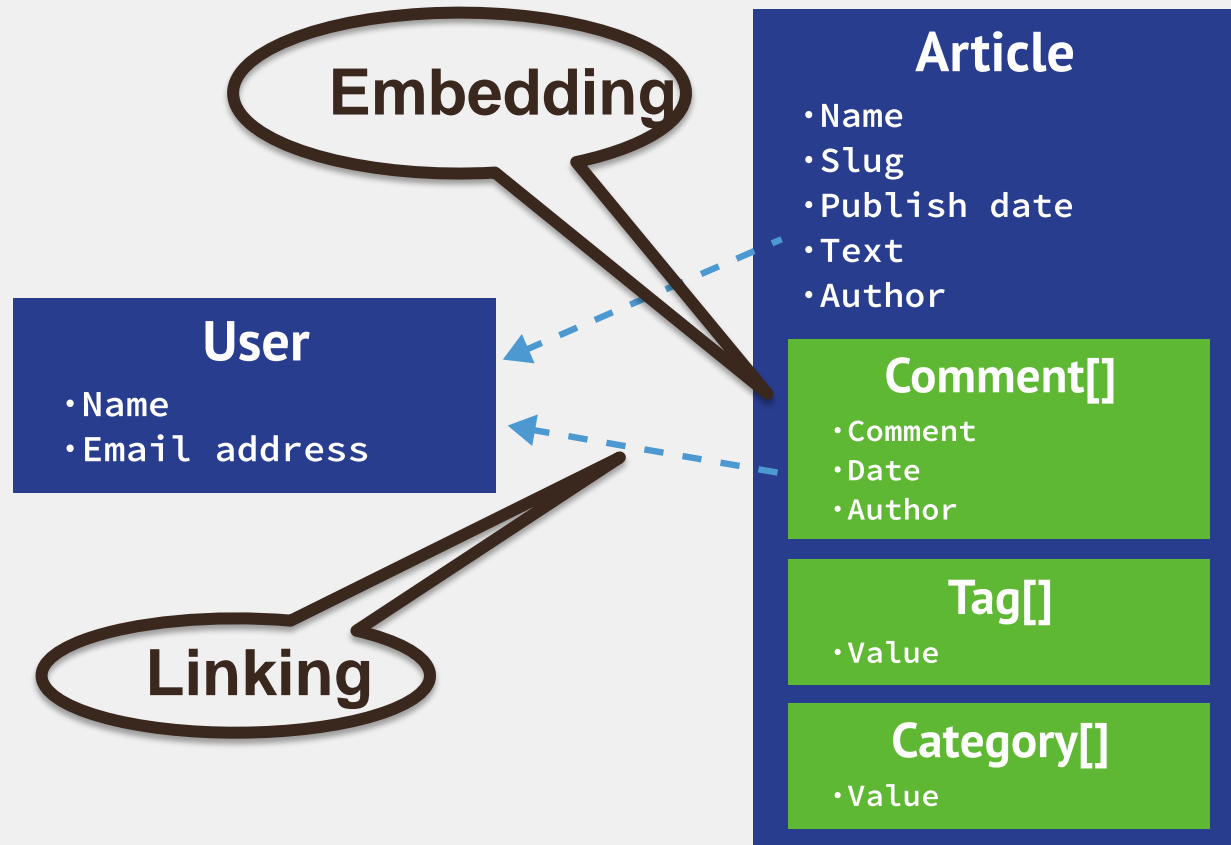
Normalized Schema (RDBMS)



De-Normalized Schema (MongoDB)



De-Normalized Schema (MongoDB)



De-Normalized Schema (MongoDB)

Relational Schema Design

Focuses on data storage

Document Schema Design

Focus on data use

Schema Design Considerations

- How do we manipulate the data?
 - Dynamic Ad-Hoc Queries
 - Atomic Updates
 - Map Reduce
- What are the access patterns of the application?
 - Read/Write Ratio
 - Types of Queries / Updates
 - Data life-cycle and growth rate

Data Manipulation

- Conditional Query Operators
 - **Scalar:** \$ne, \$mod, \$exists, \$type, \$lt, \$lte, \$gt, \$gte, \$ne
 - **Vector:** \$in, \$nin, \$all, \$size
- Atomic Update Operators
 - **Scalar:** \$inc, \$set, \$unset
 - **Vector:** \$push, \$pop, \$pull, \$pushAll, \$pullAll, \$addToSet

Data Access

- Flexible Schemas
- Ability to embed complex data structures
- Secondary Indexes
- Multi-Key Indexes
- Aggregation Framework
 - \$project, \$match, \$limit, \$skip, \$sort, \$group, \$unwind
- No Joins

Schema Design by Example

Library Management Application

- Patrons
- Books
- Authors
- Publishers

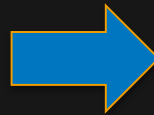
An Example

One to One Relations

Modeling Patrons

```
patron = {  
  _id: "joe"  
  name: "Joe Bookreader"  
}
```

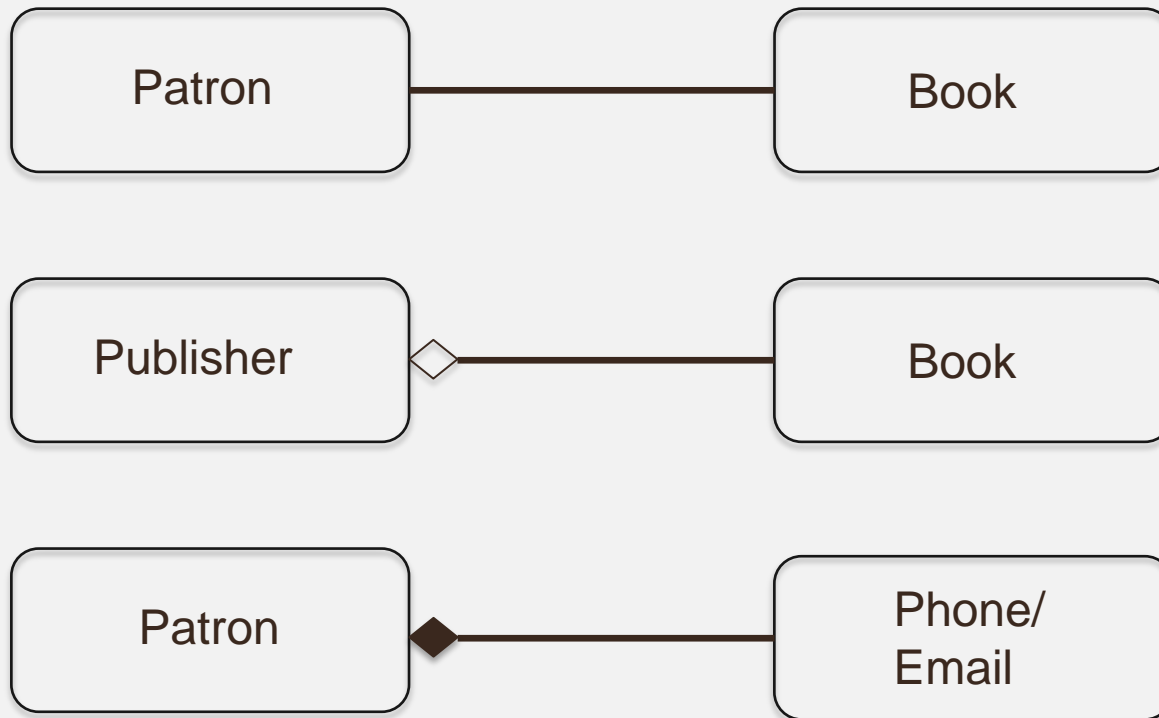
```
address = {  
  patron_id = "joe",  
  street: "123 Fake St. ",  
  city: "Faketon",  
  state: "MA",  
  zip: 12345  
}
```



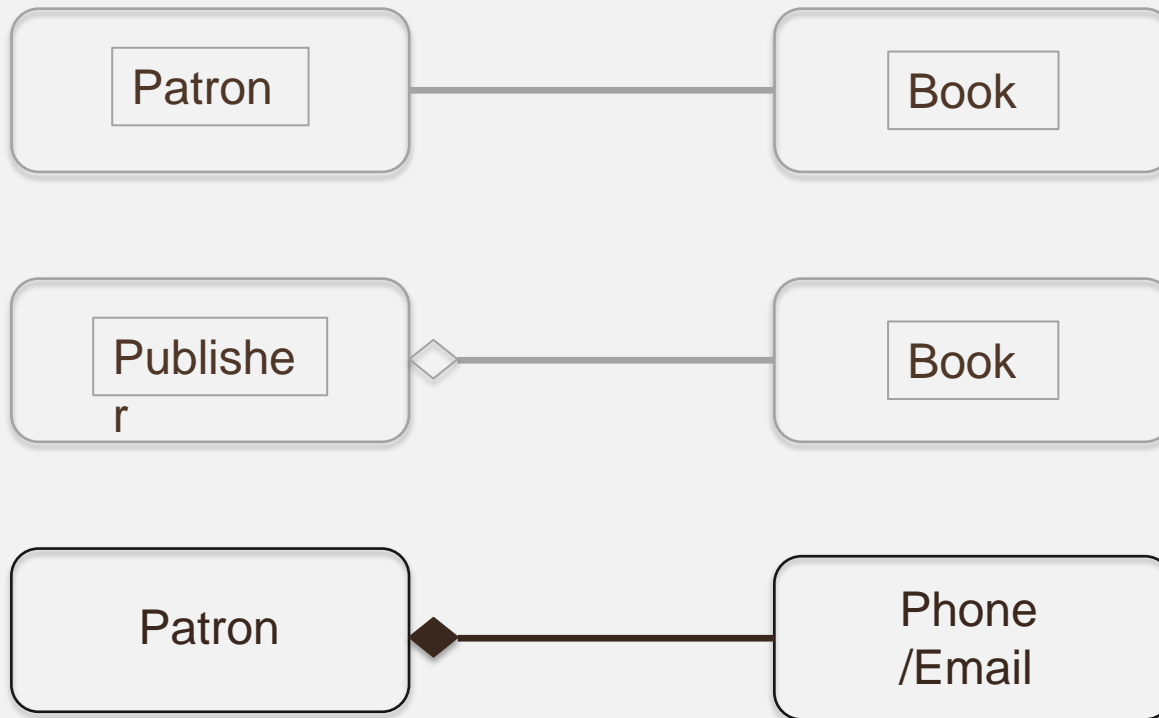
```
patron = {  
  _id: "joe"  
  name: "Joe Bookreader",  
  address: {  
    street: "123 Fake St. ",  
    city: "Faketon",  
    state: "MA",  
    zip: 12345  
  }  
}
```

An Example

One To Many Relations



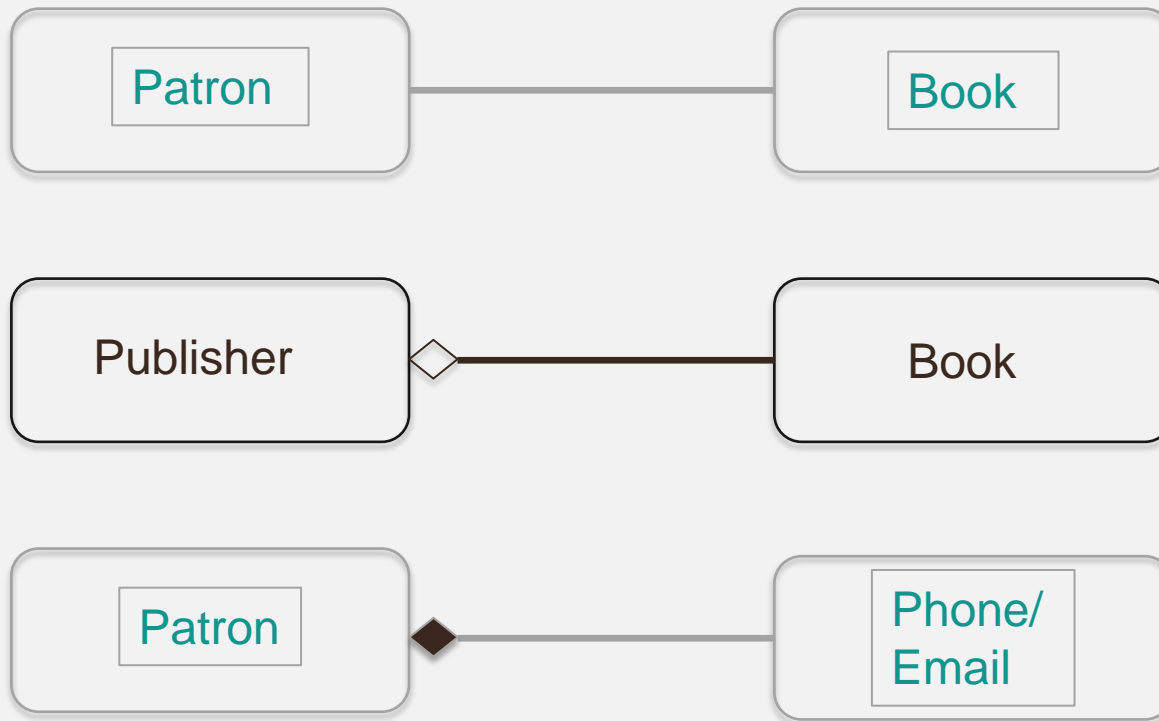
One to Many Relations



Strong ownership/lifecycle

Modeling Patrons

```
patron = {  
  _id: "joe"  
  name: "Joe Bookreader",  
  join_date: ISODate("2011-10-15"),  
  addresses: {street: "123 Fake St.", city: "Faketon", state: "MA", ...},  
  email: ["email@domain1.com", "email@domain2.com"],  
  phone:[  
    {phone: "888-111-2222", type: "home"},  
    {phone: "888-111-3333", type: "cell"}  
  ]  
}
```



Containment

Modeling Books – Embedded Publisher

```
book = {  
  title: "MongoDB: The Definitive Guide",  
  authors: [ "Kristina Chodorow", "Mike Dirolf" ],  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English",  
  
  publisher: {  
    name: "O'Reilly Media",  
    founded: "1980",  
    location: "CA"  
  }  
}
```


Modeling Books – Embedded Publisher

What is hard?

- **All publishers in the system?**
 - Distinct query
- **Publisher moved location**
 - Need to update every book published

Publisher is its own entity

```
publisher = {  
  name: "O'Reilly Media",  
  founded: "1980",  
  location: "CA"  
}
```

```
book = {  
  title: "MongoDB: The Definitive Guide",  
  authors: [ "Kristina Chodorow", "Mike Dirolf" ],  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English"  
}
```

Publisher _id as a reference

```
publisher = {  
  _id: "oreilly",  
  name: "O'Reilly Media",  
  founded: "1980",  
  location: "CA"  
}
```

```
book = {  
  title: "MongoDB: The Definitive Guide",  
  authors: [ "Kristina Chodorow", "Mike Dirolf" ],  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English",  
  publisher_id: "oreilly"  
}
```

Publisher _id as a reference

Give me books published by publisher_id.

One query

```
> db.book.find({publisher_id:"oreilly"})
```

Give me name of the publisher of "MongoDB: The Definitive Guide"

Two queries

```
> publisher_id = db.book.find({title:"MongoDB: ...."})
```

```
> db.publisher.find({_id:publisher_id})
```

Give me all books published by publishers in California

Two queries

```
> publisher_ids = db.publisher.find({location: "CA"}, {_id:1})
```

```
> db.book.find({publisher_id:{$in:publisher_ids}})
```

Book _id as a reference (embedded)

```
publisher = {  
  _id: "oreilly",  
  name: "O'Reilly Media",  
  founded: "1980",  
  location: "CA"  
  books: [ "123456789", ... ]  
}
```

```
book = {  
  _id: "123456789",  
  title: "MongoDB: The Definitive Guide",  
  authors: [ "Kristina Chodorow", "Mike Dirolf" ]  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English"  
}
```

Book _id as a reference (embedded)

Give me books published by publisher_id.

Two queries

```
> book_ids = db.publisher.find({_id:"oreilly"},{_id:0,books:1})  
> db.book.find({_id:{$in:book_ids}})
```

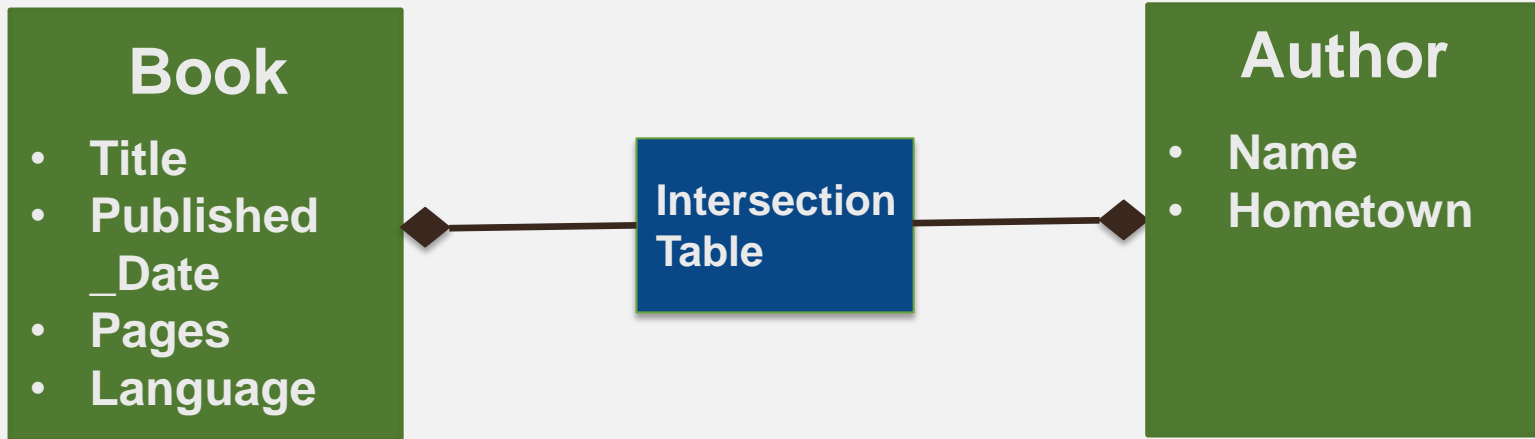
Relation stored on both sides

```
publisher = {  
  _id: "oreilly",  
  name: "O'Reilly Media",  
  founded: "1980",  
  location: "CA"  
  books: [ "123456789", ... ]  
}
```

```
book = {  
  _id: "123456789",  
  title: "MongoDB: The Definitive Guide",  
  authors: [ "Kristina Chodorow", "Mike Dirolf" ]  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English",  
  publisher id: "oreilly"  
}  
}
```

An Example

Many to Many Relations



Author

- Id
- Name
- Hometown

Relational Approach

Relation stored on book end

```
book = {  
  title: "MongoDB: The Definitive Guide",  
  authors = [  
    { _id: "kchodorow", name: "Kristina Chodorow" },  
    { _id: "mdirolf", name: "Mike Dirolf" }  
  ]  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English"  
}
```

```
author = {  
  _id: "kchodorow",  
  name: "Kristina Chodorow",  
  hometown: "New York"  
}
```

Relation stored on author end

```
book = {  
  _id: 123456789  
  title: "MongoDB: The Definitive Guide",  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English"  
}
```

```
author = {  
  _id: "kchodorow",  
  name: "Kristina Chodorow",  
  hometown: "Cincinnati",  
  books: [ {book_id: 123456789, title : "MongoDB: The Definitive Guide" } ]  
}
```

Relation stored on both sides

```
book = {  
  _id: 123456789  
  title: "MongoDB: The Definitive Guide",  
  authors = [ "kchodorow", "mdirolf" ],  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English"  
}
```

```
author = {  
  _id: "kchodorow",  
  name: "Kristina Chodorow",  
  hometown: "Cincinnati",  
  books: [ 123456789, ... ]  
}
```

Referencing vs. Embedding

- Embedding
 - Great for read performance
 - One seek to load entire object
 - One roundtrip to database
 - Writes can be slow
 - Maintaining data integrity
- Reference
 - More flexibility
 - Data integrity is maintained
 - Work is done during reads

An Example

Trees

Parent Links

```
book = {  
  title: "MongoDB: The Definitive Guide",  
  authors: [ "Kristina Chodorow", "Mike Dirolf" ],  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English",  
  category: "MongoDB"  
}  
  
category = { _id: MongoDB, parent: Databases }  
category = { _id: Databases, parent: Programming }
```

Easy to query by parent category
Hard to find in subcategories

Array of Ancestors

```
book = {  
  title: "MongoDB: The Definitive Guide",  
  authors: [ "Kristina Chodorow", "Mike Dirolf" ],  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English",  
  parent: "MongoDB",  
  categories: ["MongoDB", "Databases", "Programming" ]  
}
```

```
book = {  
  title: "MySQL: The Definitive Guide",  
  authors: [ "Michael Kofler" ],  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English",  
  parent: "MySQL",  
  ancestors: ["MySQL", "Databases", "Programming" ]  
}
```


Ancestors as path

```
book = {  
  title: "MongoDB: The Definitive Guide",  
  authors: [ "Kristina Chodorow", "Mike Dirolf" ],  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English",  
  categories: "MongoDB/Databases/Programming"  
}
```

```
book = {  
  title: "MySQL: The Definitive Guide",  
  authors: [ "Michael Kofler" ],  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English",  
  parent: "MySQL",  
  ancestors: "MySQL/Databases/Programming"  
}
```

Other common patterns

Inheritance hierarchy

- RDBMS
 - Multiple Tables vs Single Table

Shapes table

id	type	Area	radius	length	Width
1	circle	3.14	1		
2	square	4		2	
3	rect	10		5	2

- Sparse data
- Is missing value not required or an error?

Single table inheritance - RDBMS

Single collection (table) inheritance

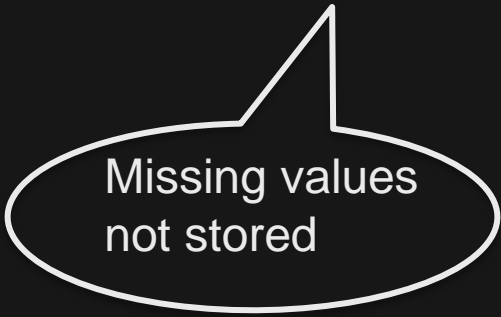
- MongoDB

```
> db.shapes.find()
```

```
{ _id : 1, type: "circle", area : 3.14, radius : 1 }
```

```
{ _id : 2, type: "square", area : 4, length : 2 }
```

```
{ _id : 3, type: "rect", area : 10, length : 5, width : 2 }
```



Missing values
not stored

Summary

- Schema design is different in MongoDB
- Basic data design principals stay the same
- Focus on how application accesses/manipulates data
- Rapidly evolve schema to meet your requirements
- Enjoy your new freedom, use it wisely 😊



#codemash

Coming Next: Part 2.2 : Indexing

Sridhar Nanjundeswaran

Engineer, MongoDB Inc.

@snanjund